

# EE-222: Microprocessor Systems

## RISC vs. CISC Architectures

Instructor: Dr. Arbab Latif

# Below Your Program: Levels of Abstraction

Higher-Level Language  
Program (e.g. C)

```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

*Compiler*

Assembly Language  
Program (e.g. RISCv)

```
lw  x5, 0(x2)  
lw  x6, 4(x2)  
sw  x6, 0(x2)
```

**EVERY INSTRUCTION IS DIRECTLY  
IMPLEMENTED IN THE HARDWARE**

Machine  
Program (RISCv)

1000  
0110

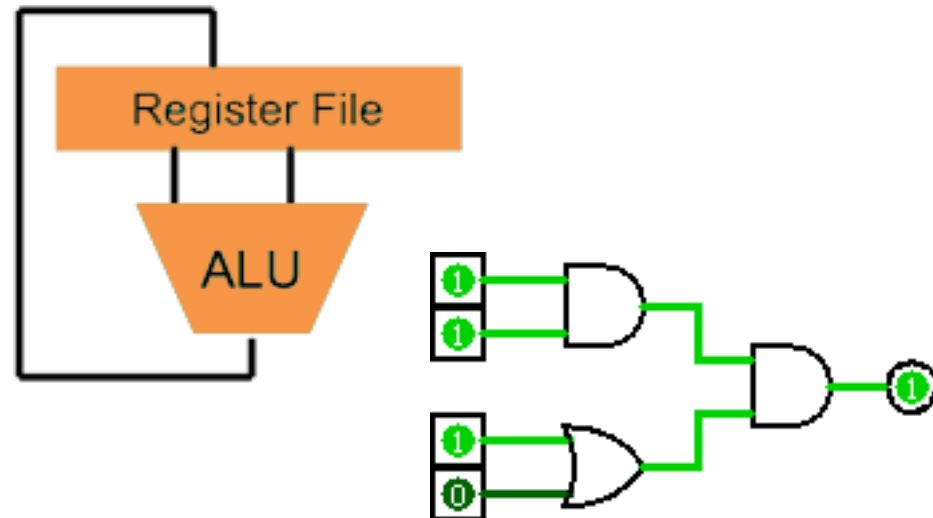
1100 0110 1010 1111 0101 1000 0000 1001  
0101 1000 0000 1001 1100 0110 1010 1111

*Machine  
Interpretation*

Hardware Architecture Description  
(e.g. block diagrams)

*Architecture  
Implementation*

Logic Circuit Description  
(Circuit Schematic Diagrams)



# Think about It!

- Every Instruction is Directly Implemented in the Hardware:
  - Think in terms of CISC vs. RISC

CISC	RISC
<b>MULT</b> 2:3, 5:2	<b>LOAD</b> A, 2:3 <b>LOAD</b> B, 5:2 <b>PROD</b> A, B <b>STORE</b> 2:3, A

- CISC:
  - Variable size instructions -> Complex instruction decoder -> More transistors
  - Multi-cycle instructions
- RISC:
  - Simplified instructions -> simple instruction decoder
  - Single-cycle instructions (mostly)

# Changing the Architecture

## CISC vs. RISC

- Early trend:
  - Add more and more instructions to do elaborate operations – *Complex Instruction Set Computing* (CISC)
    - difficult to learn and comprehend language
    - super-complicated (slow?) hardware
- Later on:
  - Opposite philosophy later began to dominate: *Reduced Instruction Set Computing* (RISC)
    - **Simpler (and smaller) instruction set makes it easier to build fast hardware**
    - Let software do the complicated operations by composing simpler ones

# RISC Architecture

- Feature 1: RISC processors have a fixed instruction size.
  - It makes the task of instruction decoder easier.
    - In AVR the instructions are 2 or 4 bytes.
  - In CISC processors instructions have different lengths
    - E.g. in 8051
      - CLR C ; a 1-byte instruction
      - ADD A, #20H ; a 2-byte instruction
      - LJMP HERE; a 3-byte instruction

# RISC architecture


- Feature 2: Reduce the number of instructions
  - Pros: Reduces the number of used transistors
  - Cons:
    - Can make the assembly programming more difficult
    - Can lead to using more memory

Feature 3: more than 95% of instructions are executed in 1 machine cycle

# What is Instruction Set Architecture (ISA)?

# The Instruction Set Architecture

- The ISA is the **interface between** what the **software** commands and what the **hardware** carries out
  - i.e the vocabulary of commands understood by a given architecture
  
- The ISA specifies
  - The **memory organization**
    - Instruction address space (ATmega16a:  $2^{13}$ )
    - Word- or Byte-addressable
  
  - The **register set**
    - 32x8 GPR registers in AVR
  
  - The **instruction set**
    - Opcodes
    - Data types
    - Addressing modes
    - Semantics of instructions





# Microarchitecture

- An **implementation** of the ISA:
  - Under specific design constraints and goals
- Anything done in hardware without exposure to software
  - Pipelining
  - In-order versus out-of-order instruction execution
  - Memory access scheduling policy
  - Speculative execution
  - Superscalar processing (multiple instruction issue?)
  - Clock gating
  - Caching? Levels, size, associativity, replacement policy
  - Prefetching?
  - Voltage/frequency scaling?
  - Error correction?

# Wrap-up: ISA vs. Microarchitecture

- ISA
  - Agreed upon interface between software and hardware
    - SW/compiler assumes, HW promises
  - What the software writer needs to know to write and debug system/user programs
- Microarchitecture
  - Specific implementation of an ISA
  - Not visible to the software
- Microprocessor
  - **ISA, uarch**, circuits
  - “Architecture” = ISA + microarchitecture

Problem
Algorithm
Program
ISA
Microarchitecture
Circuits
Electrons

# Introduction to RISC-V

# Money Makes the Mare Go!

## Choose a Licensing Plan That's Right for Your Business



### Arm DesignStart

Access a select mix of IP products and tools

[Learn More >](#)



### Arm Flexible Access

Access IP products, support, tools and training

[Learn More >](#)



### Standard Licensing

Full product portfolio, support and training

[Learn More >](#)

#### Upfront Costs

\$0 for the Cortex-M0, Cortex-M1 and Cortex-M3 CPUs

\$75K for the Cortex-A5 CPU

\$75K entry package annual access fee

\$200K standard package annual access fee

Upfront license fees based on license terms

#### Licensing


No additional license fee for the Cortex-M0, Cortex-M1, and Cortex-M3 CPUs  
\$50K license fee for the Cortex-A5 CPU at tape out

Per product license fees due at tape out

Upfront license fees based on license terms

What if there were free and open ISAs we could use for everything?

# Welcome to RISC-V



## x86

<b>Designer</b>	Intel, AMD
<b>Bits</b>	16-bit, 32-bit and 64-bit
<b>Introduced</b>	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
<b>Design</b>	CISC
<b>Type</b>	Register-memory
<b>Encoding</b>	Variable (1 to 15 bytes)
<b>Endianness</b>	Little

Macbooks & PCs  
(Core i3, i5, i7, M)  
x86 Instruction Set



## ARM architectures

<b>Designer</b>	ARM Holdings
<b>Bits</b>	32-bit, 64-bit
<b>Introduced</b>	1985; 31 years ago
<b>Design</b>	RISC
<b>Type</b>	Register-Register
<b>Encoding</b>	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 <a href="#">user-space compatibility</a> <sup>[1]</sup>
<b>Endianness</b>	Bi (little as default)

Smartphone-like devices  
(iPhone, iPad, Raspberry Pi)  
[ARM Instruction Set](#)



## RISC-V

<b>Designer</b>	University of California, Berkeley
<b>Bits</b>	32, 64, 128
<b>Introduced</b>	2010
<b>Version</b>	2.2
<b>Design</b>	RISC
<b>Type</b>	Load-store
<b>Encoding</b>	Variable
<b>Branching</b>	Compare-and-branch
<b>Endianness</b>	Little

Versatile and open-source  
Relatively new, designed for  
cloud computing, high-end  
phones, small embedded sys.  
[RISCV Instruction Set](#)

# What is RISC-V?

- RISC-V (pronounced "risk-five")
  - Is an open specification of an Instruction Set Architecture (ISA) based on reduced instruction set computing (RISC) 5<sup>th</sup> iteration:
    - Describes the way in which software talks to an underlying processor
  - There was RISC-I, II, III, IV before
- Most ISAs: X86, ARM, Power, MIPS, SPARC
  - Commercially protected by patents
  - Preventing practical efforts to reproduce the computer systems.
- RISC-V is open
  - Permitting any person or group to construct compatible computers
  - Use associated software



<https://riscv.org/>

# What's Different About RISC-V?

- Simple:
  - Far smaller than proprietary ISAs
  - 2500 pages for x86, ARMv8 manual vs 200 for RISC-V manual
    - 47 RISC-V instructions vs 3500+ 80x86 instructions
- Clean-slate design:
  - 25 years later, so can learn from mistakes of predecessors
  - Avoids architecture or technology-dependent features



# What's Different About RISC-V?

- Modular:
  - A mandatory **Base integer ISA**
    - **I: Integer instructions:**
      - ALU
      - Branches/jumps
      - Loads/stores
  - **Standard Extensions**
    - **M: Integer Multiplication and Division**
    - **A: Atomic Instructions**
    - **F: Single-Precision Floating-Point**
    - **D: Double-Precision Floating-Point**

# What's Different About RISC-V?

- Modular:

Name of base or extension	Functionality
RV32I	Base 32-bit integer instruction set with 32 registers
RV32E	Base 32-bit instruction set but with only 16 registers; intended for very low-end embedded applications
RV64I	Base 64-bit instruction set; all registers are 64-bits, and instructions to move 64-bit from/to the registers (LD and SD) are added
M	Adds integer multiply and divide instructions
A	Adds atomic instructions needed for concurrent processing; see Chapter 5
F	Adds single precision (32-bit) IEEE floating point, includes 32 32-bit floating point registers, instructions to load and store those registers and operate on them
D	Extends floating point to double precision, 64-bit, making the registers 64-bits, adding instructions to load, store, and operate on the registers
Q	Further extends floating point to add support for quad precision, adding 128-bit operations
L	Adds support for 64- and 128-bit decimal floating point for the IEEE standard
C	Defines a compressed version of the instruction set intended for small-memory-sized embedded applications. Defines 16-bit versions of common RV32I instructions
V	A future extension to support vector operations (see Chapter 4)
B	A future extension to support operations on bit fields
T	A future extension to support transactional memory
P	An extension to support packed SIMD instructions: see Chapter 4
RV128I	A future base instruction set providing a 128-bit address space

# What's Different About RISC-V?

- Supports specialization:
  - Vast opcode space reserved
- Community designed:
  - Base and standard extensions finished
  - Grow via optional extensions vs. incremental required features
- RISC-V Foundation extends ISA for technical reasons:
  - vs. private corporation for internal (marketing) reasons

# Class ISA is RISC-V

- RISC-V is a new free, simple, clean, extensible ISA developed at Berkeley for education and research
  - Both of the dominant ISAs (x86 and ARM) are too complex to use for teaching or research
- RISC-V Foundation manages standard [riscv.org](https://riscv.org)
- RISC-V has taken off commercially
  - Nvidia is using RISC-V in all future GPUs
  - Western Digital is using RISC-V in all future products
- Now upstream support for many tools (gcc, Linux, FreeBSD, ...)
- Govt. India selected RISC-V as national ISA



# RISC-V



# Why RISC-V

- Open and free
- Not domain-specific
- To keep things simple, flexible and extensible
- No baggage of legacy

# RISC-V ISA

# RISC-V ISA Design Principle-1

- Design Principle 1: Simplicity favours regularity
  - Regularity makes implementation simpler
  - Simplicity enables higher performance at lower cost
- e.g. All arithmetic operations have same form
  - Two sources and one destination
- add a, b, c // a gets b + c



# RISC-V ISA Design Principle-2

- Design Principle 2: Smaller is faster
  - memory is larger than no. of registers, use register operands
- e.g. Arithmetic operations use register operands and not direct memory
- most implementations have decoding the operands on the critical path so only 32 registers

# RISC-V ISA Design Principle-3

- Design Principle 3: Make the common case fast
  - Small constants are common
  - Immediate operand avoids a load instruction
- support for immediate operands,
- e.g. `addi x22, x22, 4`

# RISC-V ISA Design Principle-4

- Design Principle 4: Good design demands good compromises
  - Different formats complicate decoding, but allow 32-bit instructions uniformly
  - Keep formats as similar as possible
- e.g. R-format, I-format, S-format

# RISC-V ISA

# RISC-V – How Many Registers?

- RISC-V has a  $32 \times 64$ -bit register file [RV64]
  - RV32 has  $32 \times 32$  register file
  - Use for frequently accessed data
  - 64-bit data is called a “doubleword”
    - 32 x 64-bit general purpose registers x0 to x30
  - 32-bit data is called a “word”

REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# RISC-V – How Many Registers?

## REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# RISC-V Processor State

- Program counter (PC)
- 32 32/64-bit integer registers (x0-x31)
  - x0 always contains a 0
  - x1 to hold the return address on a call.
- 32 floating-point (FP) registers (f0-f31)
  - Each can contain a single- or double-precision FP value (32-bit or 64-bit IEEE FP)

REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# Four Core RISC-V Instruction Formats

Instruction format	Primary use	rd	rs1	rs2	Immediate
R-type	Register-register ALU instructions	Destination	First source	Second source	
I-type	ALU immediates Load	Destination	First source base register		Value displacement
S-type	Store Compare and branch		Base register first source	Data source to store second source	Displacement offset
U-type	Jump and link Jump and link register	Register destination for return PC	Target address for jump and link register		Target address for jump and link

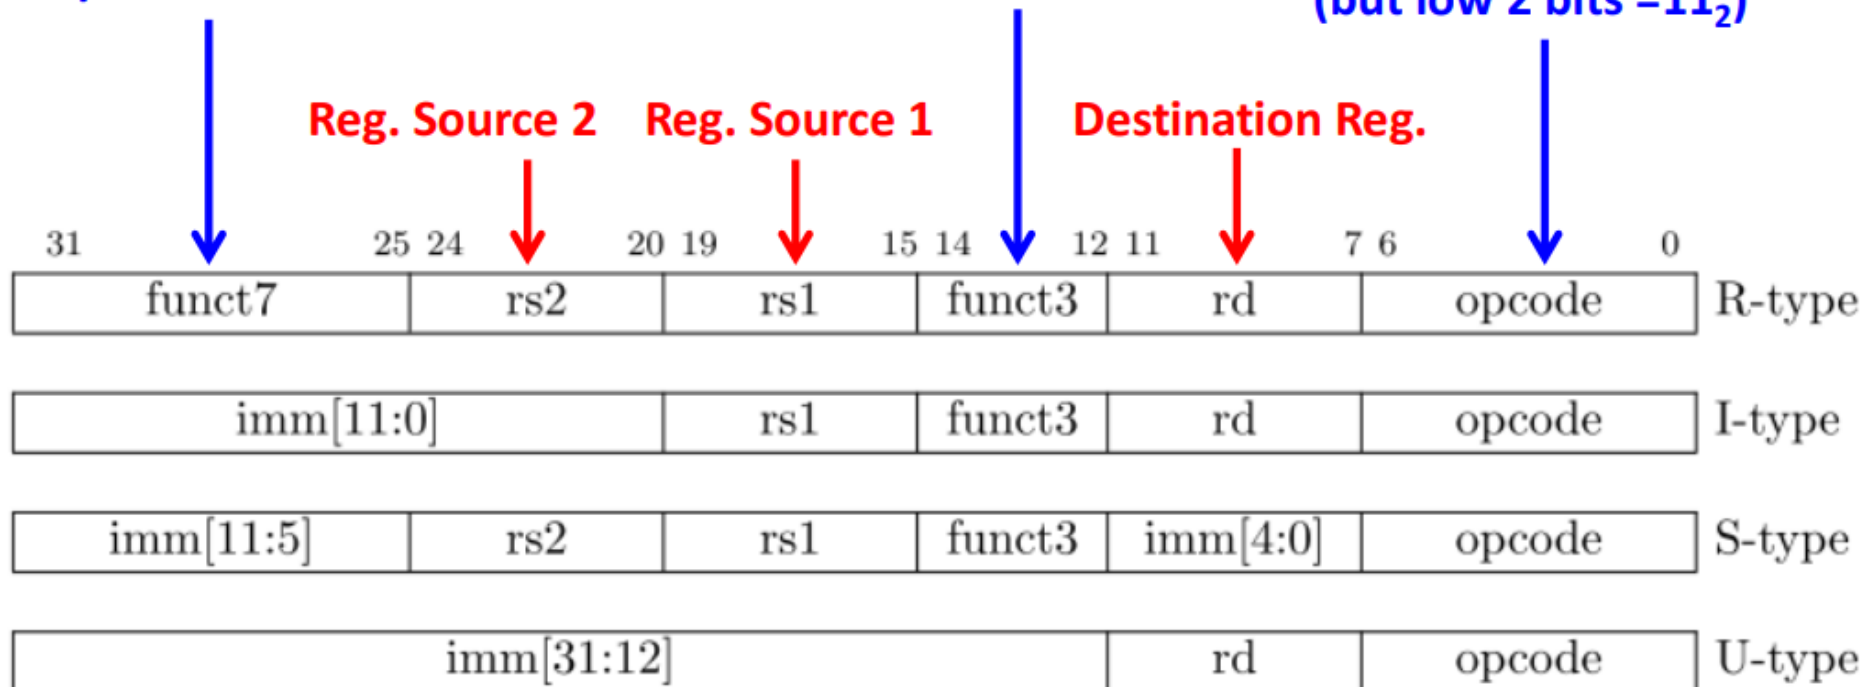


# Four Core RISC-V Instruction Formats

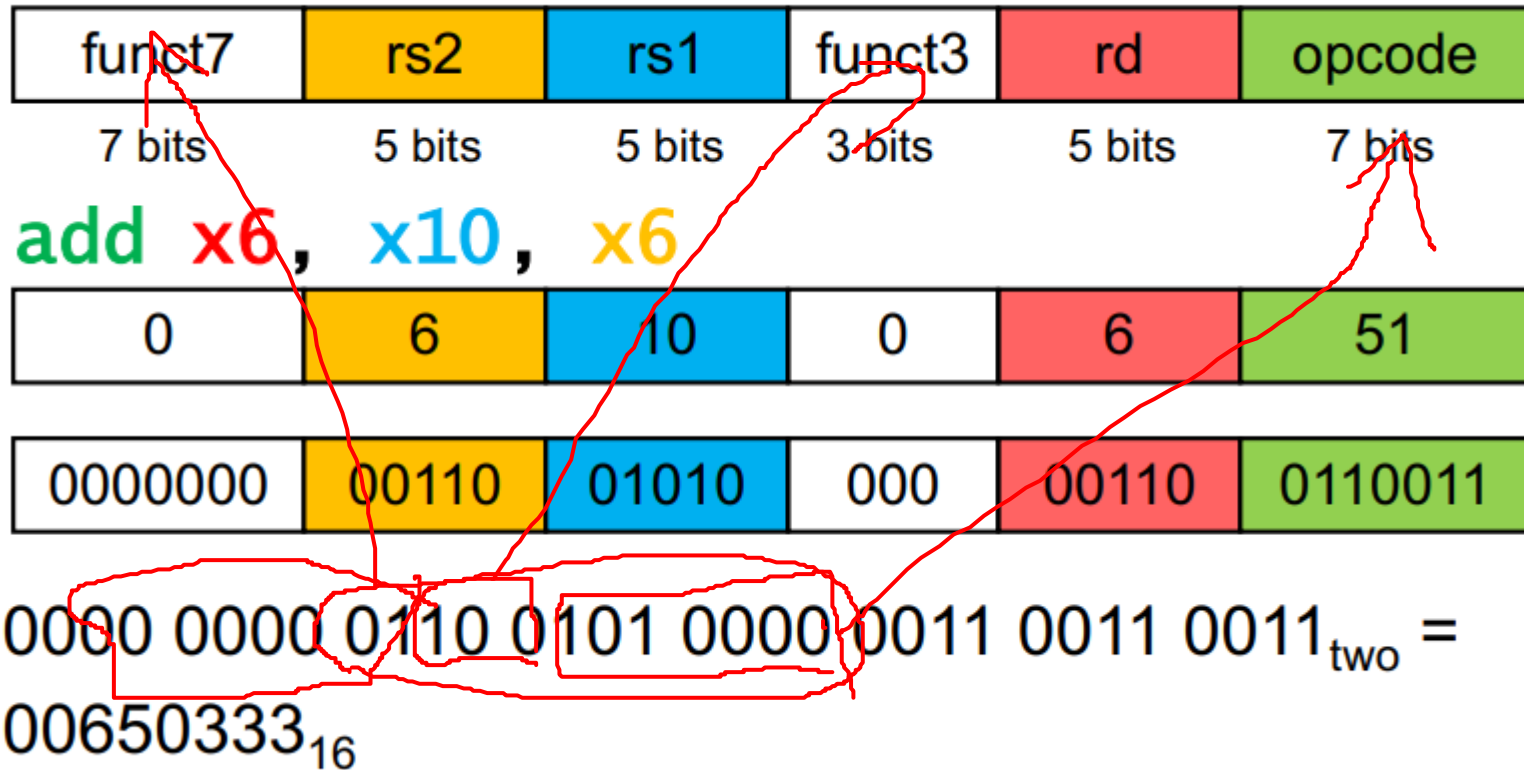
<https://github.com/riscv/riscv-opcodes/blob/master/opcodes>

Additional opcode  
bits/immediate

Additional opcode bits  
7-bit opcode field  
(but low 2 bits =  $11_2$ )



# R-Format Encoding Example



# Specifications and Software

## From riscv.org and github.com/riscv

- Specification from RISC-V website
  - <https://riscv.org/specifications/>
- RISC-V software includes
  - GNU Compiler Collection (GCC) toolchain (with GDB, the debugger)
    - <https://github.com/riscv/riscv-tools>
  - LLVM toolchain
  - A simulator ("Spike")
    - <https://github.com/riscv/riscv-isa-sim>
  - Standard simulator QEMU
    - <https://github.com/riscv/riscv-qemu>
- Operating systems support exists for Linux
  - <https://github.com/riscv/riscv-linux>
- A JavaScript ISA simulator to run a RISC-V Linux system on a web browser
  - <https://github.com/riscv/riscv-angel>

# Class Tools

- Venus Simulator: <https://venus.cs61c.org/>
- 5-stage processor pipeline simulator and assembly code editor
  - <https://github.com/mortbopet/Ripes>

# RISC-V Implementations

- A list from
  - <https://riscv.org/risc-v-cores/>
- The Indian IIT-Madras is developing six RISC-V open-source CPU designs (SHAKTI) for six distinct usages
  - <https://shaktiproject.bitbucket.io/index.html>
- SiFive HiFive Unleashed
  - First Linux RISC-V Board
    - First shipment: June 2018
  - <https://www.sifive.com/>
  - <https://github.com/sifive/freedom>



# Recommended Reading

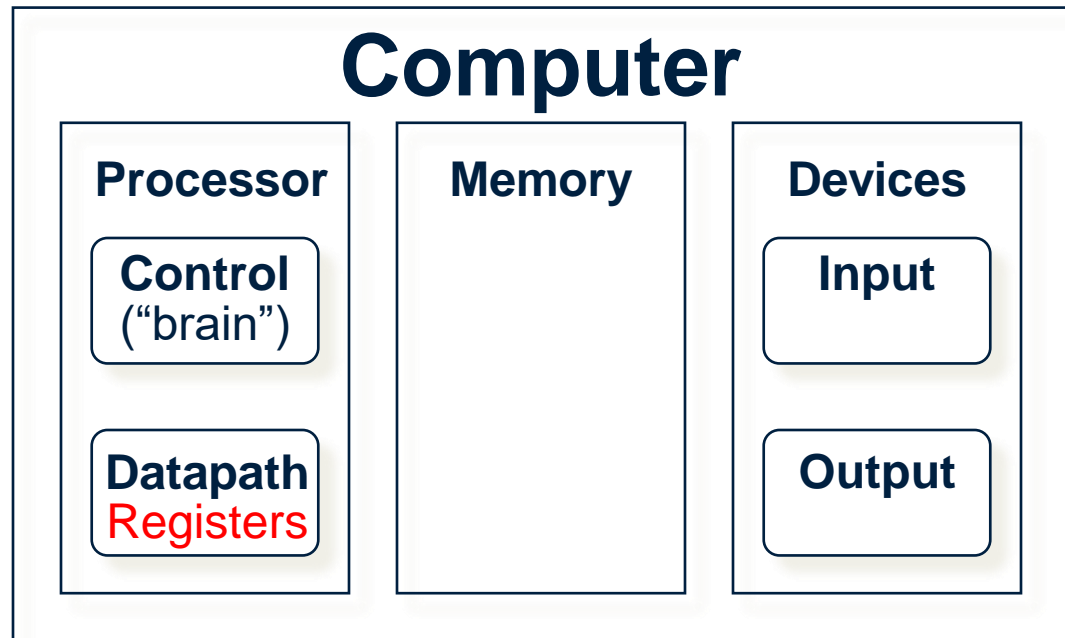
- RISC-V Edition - Computer Organization and Design: The Hardware Software Interface **(COD)** - David A. Patterson, John L. Hennessy:
  - Chapter-2

# Variables in Hardware

- Already know instructions must be directly implemented in hardware
- What about variables?!?!
  - Live in Memory
  - Relative to your processor, memory is soooooo slooooooowwwwwwwwwwwwwwwwwww
  - *predetermined* small number of fast **registers** in processor to hold variables.

# Registers in a Computer

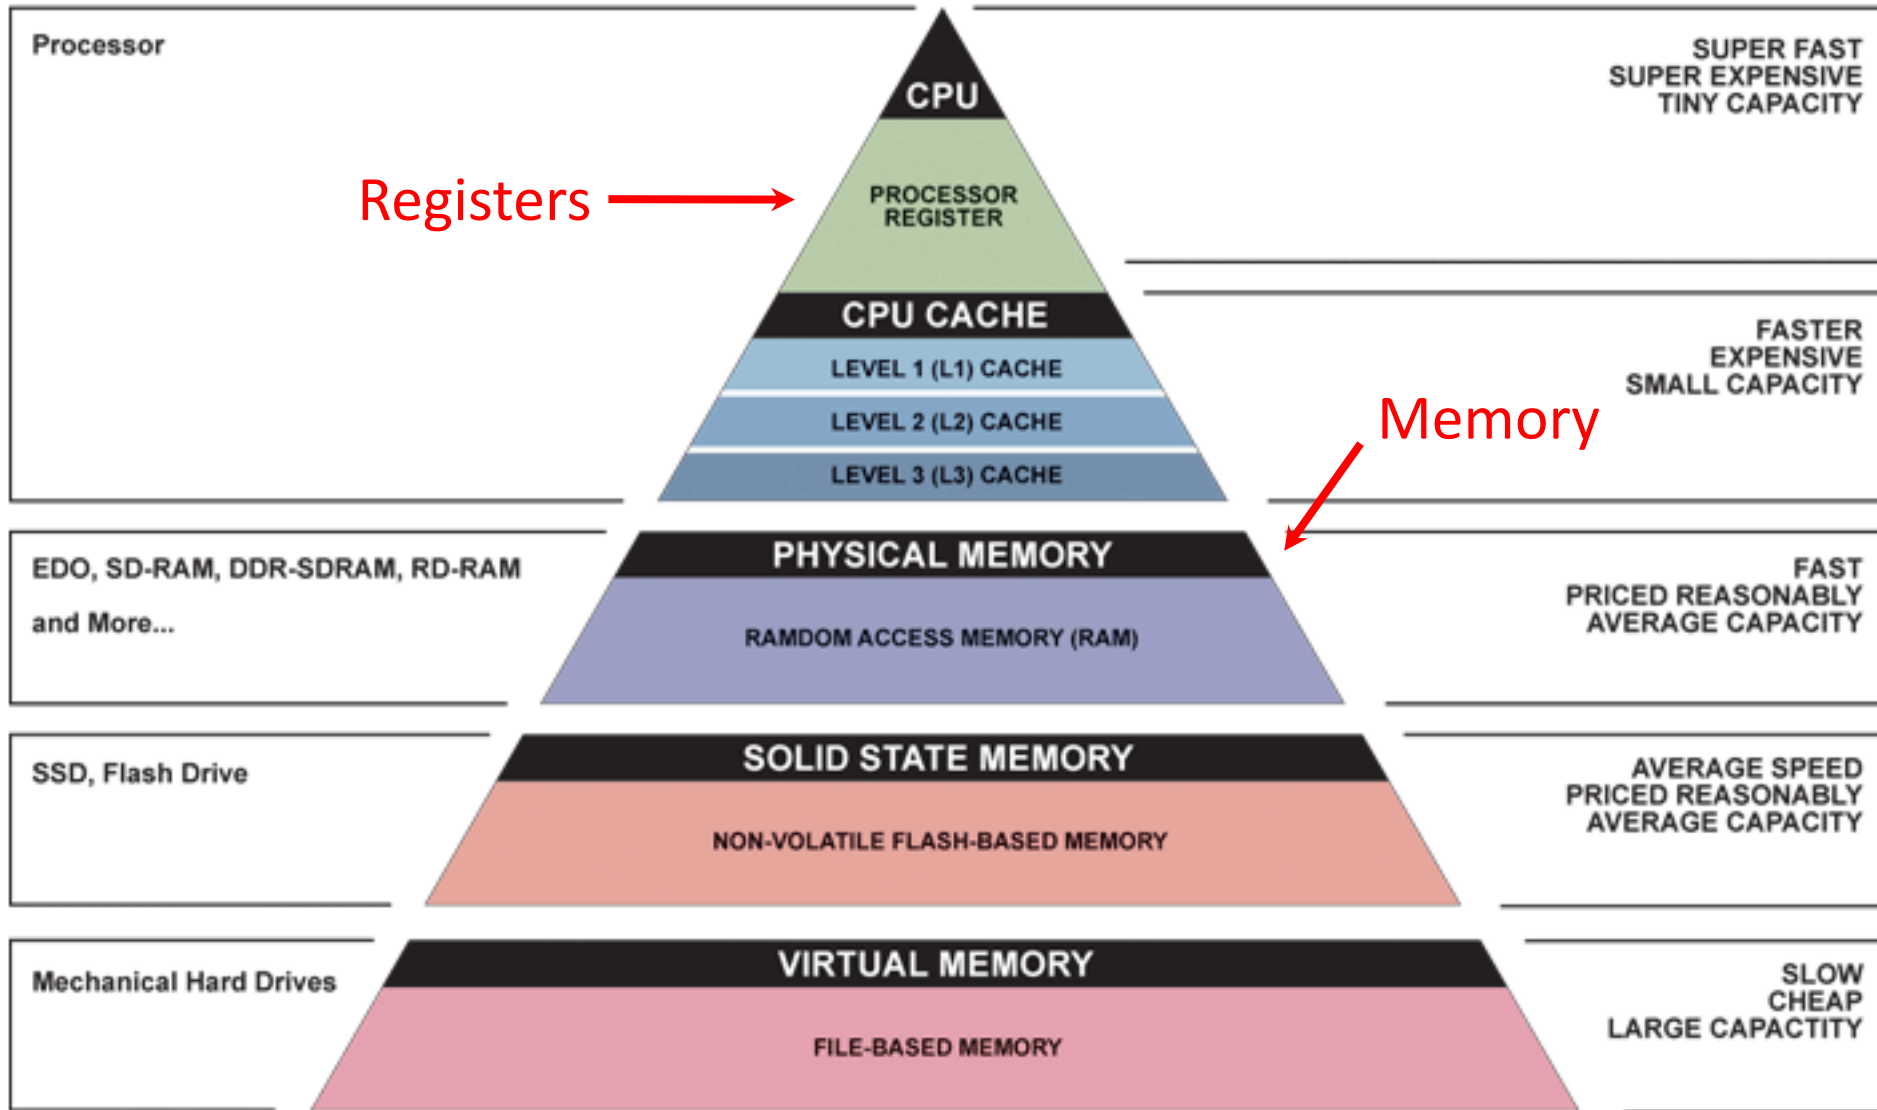
- Five components of a computer!
  - Control
  - Datapath
  - Memory
  - Input
  - Output



- Registers are part of the Datapath



# Memory Hierarchy



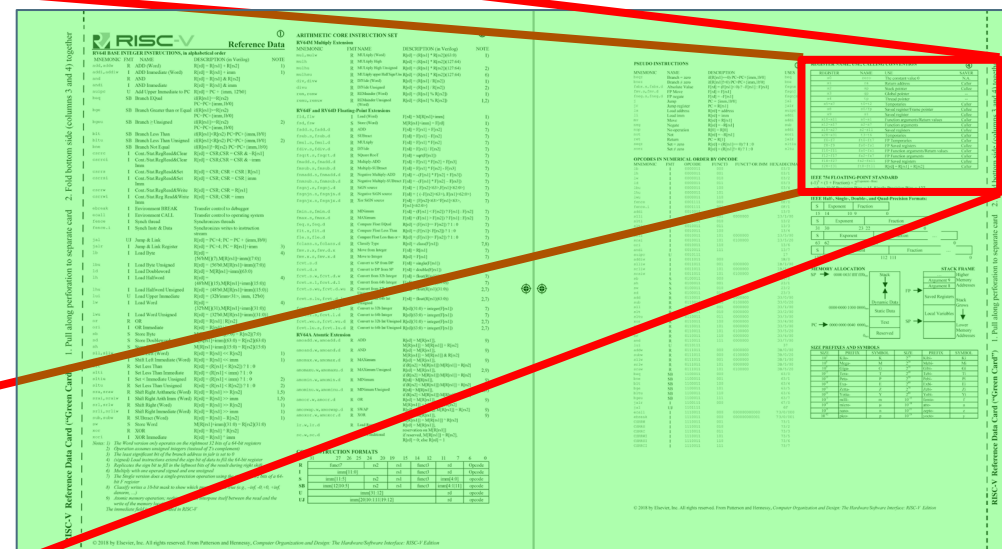
# Register vs. Memory

- What if more variables than registers?
  - Keep most frequently used in registers and move the rest to memory (called *spilling* to memory)
- Why not all variables in memory?
  - Smaller is faster: registers 100-500 times faster
  - Registers more versatile
    - In 1 arithmetic instruction: read 2 operands, perform 1 operation, and 1 write
    - In 1 data transfer instruction: 1 memory read/write, no operation

# RISC-V – How Many Registers?

- Tradeoff between speed and availability
  - more registers → can house more variables simultaneously; all registers are slower.
- RISC-V has 32 registers (x0-x31)
  - Each register is 32 bits wide and holds a word

REGISTER NAME, USE, CALLING CONVENTION		
REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$



# RISC-V – How Many Registers?

## REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# RISC-V Registers

- Register denoted by 'x' can be referenced by number (x0-x31) or name:
  - Registers that hold programmer variables:

s0-s1	↔	x8-x9
s2-s11	↔	x18-x27
  - Registers that hold temporary variables:

t0-t2	↔	x5-x7
t3-t6	↔	x28-x31
  - You'll learn about the other 13 registers later
- *Registers have no type* (C concept); the operation being performed determines how register contents are treated

# RISC-V Processor State

- Program counter (PC)
- 32 32/64-bit integer registers (x0-x31)
  - x0 always contains a 0
  - x1 to hold the return address on a call.
- 32 floating-point (FP) registers (f0-f31)
  - Each can contain a single- or double-precision FP value (32-bit or 64-bit IEEE FP)

REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# Side Note: RV64

- RISC-V has a  $32 \times 64$ -bit register file [RV64]
  - RV32 has  $32 \times 32$  register file
  - Use for frequently accessed data
  - 64-bit data is called a “doubleword”
    - 32 x 64-bit general purpose registers x0 to x30
  - 32-bit data is called a “word”

REGISTER NAME, USE, CALLING CONVENTION

REGISTER	NAME	USE
x0	zero	The constant value 0
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7	t0-t2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-x11	a0-a1	Function arguments/Return values
x12-x17	a2-a7	Function arguments
x18-x27	s2-s11	Saved registers
x28-x31	t3-t6	Temporaries
f0-f7	ft0-ft7	FP Temporaries
f8-f9	fs0-fs1	FP Saved registers
f10-f11	fa0-fa1	FP Function arguments/Return values
f12-f17	fa2-fa7	FP Function arguments
f18-f27	fs2-fs11	FP Saved registers
f28-f31	ft8-ft11	$R[rd] = R[rs1] + R[rs2]$

# Registers: Summary

- In high-level languages, number of variables limited only by available memory
- ISAs have a fixed, small number of operands called **registers**
  - Special locations built directly into hardware
  - **Benefit:** Registers are EXTREMELY FAST (faster than 1 billionth of a second)
  - **Drawback:** Operations can only be performed on these predetermined number of registers



# Getting to Know RISC-V Assembly

<http://digitalassets.lib.berkeley.edu/techreports/ucb/text/EECS-2016-1.pdf> (Architecture)



SIZE	PREFIX	SYMBOL	SIZE	PREFIX	SYMBOL
$10^0$	Kilo	K	$10^3$	Kilo	K
$10^1$	Mega	M	$10^6$	Mega	M
$10^2$	Giga	G	$10^9$	Giga	G
$10^3$	Tera	T	$10^{12}$	Tera	T
$10^4$	Peta	P	$10^{15}$	Peta	P
$10^5$	Exa	E	$10^{18}$	Exa	E
$10^6$	Zetta	Z	$10^{21}$	Zetta	Z
$10^7$	Yotta	Y	$10^{24}$	Yotta	Y
$10^8$	milli	m	$10^{-11}$	femto	f
$10^9$	micro	$\mu$	$10^{-12}$	atto	a
$10^{10}$	nano	n	$10^{-13}$	zepto	z
$10^{11}$	pico	p	$10^{-14}$	yocto	y

57

# RISC-V Instructions

- Instruction Syntax

`op dst, src1, src2`

- 1 operator, 3 operands

- `op` = operation name (“operator”)
- `dst` = register getting result (“destination”)
- `src1` = first register for operation (“source 1”)
- `src2` = second register for operation (“source 2”)

# RISC-V Instructions

- One operation per instruction,  
at most one instruction per line
- Assembly instructions are related to C operations  
(=, +, -, \*, /, &, |, etc.)
  - Must be, since C code decomposes into assembly!
  - A single line of C may break up into several lines of RISC-V

# Basic Arithmetic Instructions

# RISCV Instructions Example

- Your very first instructions!  
(assume here that the variables `a`, `b`, and `c` are assigned to registers `s1`, `s2`, and `s3`, respectively)
- Integer Addition (`add`)
  - C:  $a = b + c$
  - RISCV: `add s1, s2, s3`
- Integer Subtraction (`sub`)
  - C:  $a = b - c$
  - RISCV: `sub s1, s2, s3`

# RISCV Instructions Example

- Suppose  $a \rightarrow s0$ ,  $b \rightarrow s1$ ,  $c \rightarrow s2$ ,  $d \rightarrow s3$  and  $e \rightarrow s4$ . Convert the following C statement to RISCV:

$a = (b + c) - (d + e);$

```
add t1, s3, s4
add t2, s1, s2
sub s0, t2, t1
```

Ordering of  
instructions matters  
(must follow order  
of operations)

Utilize temporary registers

# Comments



# Comments in RISC-V

- Comments in RISC-V follow hash mark (#) until the end of line
  - Improves readability and helps you keep track of variables/registers!

```
add t1, s3, s4      # t1=d+e
add t2, s1, s2      # t2=b+c
sub s0, t2, t1      # a= (b+c) - (d+e)
```

x0 [ZERO]

# The Zero Register

- Zero appears so often in code and is so useful that it has its own register!
- Register zero (`x0` or `zero`) always has the value 0 and cannot be changed!
  - i.e. any instruction with `x0` as `dst` has no effect
- Example Uses:
  - `add s3, x0, x0 # c=0`
  - `add s1, s2, x0 # a=b`

# Immediates

# Immediates

- Numerical constants are called **immediates**
- **Why Immediate Instructions?**
  - Immediate operand avoids a Load Instruction
- Separate instruction syntax for immediates:  
`opi dst, src, imm`
  - Operation names end with 'i', replace 2<sup>nd</sup> source register with an immediate
- Example Uses:
  - `addi s1, s2, 5 # a=b+5`
  - `addi s3, s3, 1 # c++`
- **Why no subi instruction?**
  - `subi dst, src, imm = addi dst, src, -imm.`

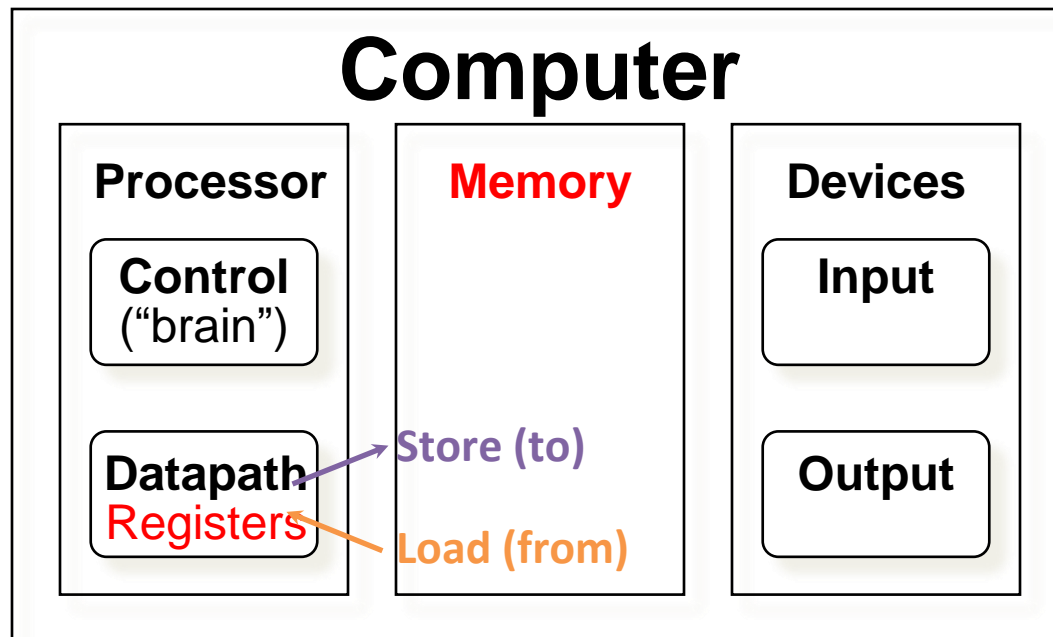
# Recommended Reading

- RISC-V Edition - Computer Organization and Design\_ The Hardware Software Interface - David A. Patterson, John L. Hennessy:
  - Chapter-2

# Data Transfer Instructions

# Data Transfer Instructions

- Data Transfer instructions are between registers (Datapath) and Memory
  - Allow us to fetch and store operands in memory





# Data Transfer

- C variables map onto registers;  
What about large data structures like arrays?
  - Don't forget *memory*, our one-dimensional array indexed by addresses starting at 0
- RISC-V instructions only operate on registers!
- Specialized data transfer instructions move data between registers and memory
  - Store (sw): register TO memory
  - Load (lw): register FROM memory

# Data Transfer


- Instruction syntax for data transfer:

`memop reg, off(bAddr)`

- `memop` = operation name (“operator”)
  - `reg` = register for operation source or destination
  - `bAddr` = register with pointer to memory (“base address”)
  - `off` = address offset (immediate) in bytes (“offset”)
- Accesses memory at address `bAddr+off`
- **Reminder:** A register holds a word of raw data (no type) – make sure to use a register (and offset) that point to a valid memory address

# Memory

- Memory is just a large, single-dimensional array, with the address acting as the index to that array, starting at 0.



# Memory is Byte-Addressed

- What was the smallest data type we saw in C?

- A char, which was a *byte* (8 bits)
- Everything in multiples of 8 bits (e.g. 1 word = 4 bytes)

Assume here addr of lowest byte in word is addr of word




...	...	...	...
<u>12</u>	13	14	15
<u>8</u>	9	10	11
<u>4</u>	5	6	7
<u>0</u>	1	2	3

- Memory addresses are indexed by *bytes*, not words
- Word addresses are 4 bytes apart
  - Word addr is same as left-most byte
  - Addrs must be multiples of 4 to be “word-aligned”
- Pointer arithmetic not done for you in assembly
  - Must take data size into account yourself

# Double-word Memory Address

- Double word addresses are 8 bytes



# Data Transfer Instructions

- **Load Word** (`lw`)
  - Takes data at address `bAddr+off` FROM memory and places it into `reg`
- **Store Word** (`sw`)
  - Takes data in `reg` and stores it TO memory at address `bAddr+off`
- **Example Usage:**

```
# addr of int A[] -> s3, a -> s2
lw    t0, 12(s3) # $t0=A[3]
add   t0, s2, t0 # $t0=A[3]+a
sw    t0, 40(s3) # A[10]=A[3]+a
```

**Remember! Pointer arithmetic not done for you in assembly**  
—Must take data size into account yourself

# Recommended Reading

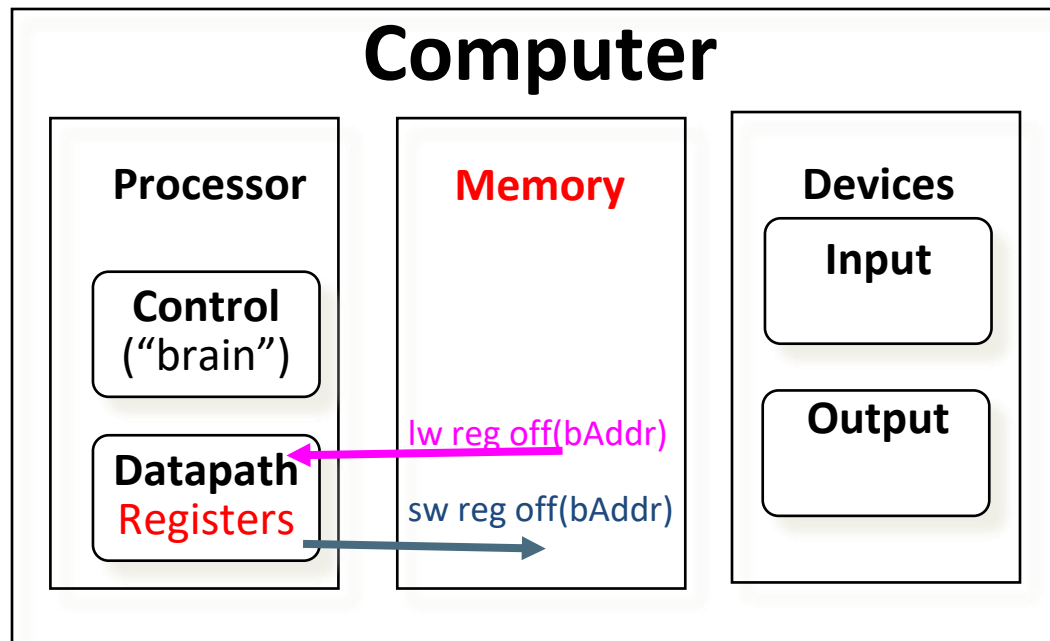
- RISC-V Edition - Computer Organization and Design\_ The Hardware Software Interface - David A. Patterson, John L. Hennessy:
  - Chapter-2

## Data Transfer Instructions Cont....



# Memory and Variable Size

- So Far:
  - lw reg, off(bAddr)
  - sw reg, off(bAddr)



- What about characters (1B) and shorts (sometimes 2B), etc?
  - Want to be able to use interact with memory values smaller than a word.

# Trading Bytes with Memory (2 approaches)

- Method 1: Move words in and out of memory using bitmasking and shifting

```
lw    s0, 0(s1)
```

```
andi  s0, s0, 0xFF # lowest byte
```

- Method 2: Load/store byte instructions

```
lb    s1, 1(s0)
```

```
sb    s1, 0(s0)
```


\* (s0) = 0x00000180

00	00	01	80
----	----	----	----

# Endianness


- **Big Endian:** Most-significant byte at least address of word
  - word address = address of most significant byte
- **Little Endian:** Least-significant byte at least address of word
  - word address = address of least significant byte

$$*(s0) = 0x00000180$$




- **RISC-V is Little Endian**

# Byte Instructions



- `lb/sb` utilize the **least significant byte of the register**
  - On `sb`, upper 24 bits are ignored
  - On `lb`, upper 24 bits are filled by sign-extension
- For example, let `* (s0) = 0x00000180`:
  - `lb s1, 1(s0)`    # `s1=0x00000001`
  - `lb s2, 0(s0)`    # `s2=0xFFFFFFFF80`
  - `sb s2, 2(s0)`    # `* (s0)=0x00800180`

# Half-Word Instructions



- `lh reg, off(bAddr)` “load half”
- `sh reg, off(bAddr)` “store half”
  - On `sh`, upper 16 bits are ignored
  - On `lh`, upper 16 bits are filled by sign-extension

## Unsigned Instructions

- `lhu reg, off(bAddr)` “load half unsigned”
- `lbu reg, off(bAddr)` “load byte unsigned”
  - On `l(b/h)u`, upper bits are filled by zero-extension
- Why no `s(h/b)u`? Why no `lwu`?

# Decision Making Instructions

# Computer Decision Making

- In C, we had *control flow*
  - Outcomes of comparative/logical statements determined which blocks of code to execute
- In RISC-V, we can't define blocks of code; all we have are **labels**
  - Defined by text followed by a colon (e.g. `main:`) and refers to the instruction that follows
  - Generate control flow by jumping to labels
  - C has these too, but they are considered bad style

# Decision Making Instructions

- **Branch If Equal** (beq)
  - `beq reg1, reg2, label`
  - If value in `reg1` = value in `reg2`, go to `label`
- **Branch If Not Equal** (bne)
  - `bne reg1, reg2, label`
  - If value in `reg1`  $\neq$  value in `reg2`, go to `label`
- **Jump** (j)
  - `j label`
  - Unconditional jump to `label`



# Breaking Down the If Else

## ■ C code:


```
if (i==j) f = g+h;  
else f = g-h;
```

- f, g, ... in x19, x20, ...

## ■ Compiled RISC-V code:

```
    bne x22, x23, Else  
    add x19, x20, x21  
    beq x0,x0,Exit // unconditional  
Else: sub x19, x20, x21  
Exit: ...
```

Assembler calculates addresses



# Breaking Down the If Else

## C Code:

```
if (i==j) {  
    a = b    /* then */  
} else {  
    a = -b   /* else */  
}
```

## In English:

- If TRUE, execute the THEN block
- If FALSE, execute the ELSE block

## RISCV (beq):

```
# i→s0, j→s1
```

```
# a→s2, b→s3
```

```
beq s0, s1, ???
```

??? — This label unnecessary

```
sub s2, x0, s3
```

```
j      end
```

```
then:
```

```
add s2, s3, x0
```

```
end:
```

# Breaking Down the If Else

## C Code:

```
if (i==j) {  
    a = b    /* then */  
} else {  
    a = -b   /* else */  
}
```

## In English:

- If TRUE, execute the THEN block
- If FALSE, execute the ELSE block

## RISCV (bne):

```
# i→s0, j→s1  
# a→s2, b→s3  
  
bne s0,s1,???  
???  
add s2, s3, x0  
j     end  
  
else:  
sub s2, x0, s3  
end:
```

# Branching on Conditions other than (Not) Equal

- **Branch Less Than (blt)**
  - `blt reg1, reg2, label`
  - If value in `reg1` < value in `reg2`, go to `label`
- **Branch Greater Than or Equal (bge)**
  - `bge reg1, reg2, label`
  - If value in `reg1` >= value in `reg2`, go to `label`
- **Example**
  - if (`a` > `b`) `a += 1`;
  - `a` in `x22`, `b` in `x23`  
`bge x23, x22, Exit`      // branch if `b` >= `a`  
`addi x22, x22, 1`

Exit:

# Signed vs. Unsigned

- Signed comparison: **blt, bge**
- Unsigned comparison: **bltu, bgeu**

- Example

- `x22 = 1111 1111 1111 1111 1111 1111 1111 1111`

- `x23 = 0000 0000 0000 0000 0000 0000 0000 0001`

- `x22 < x23 // signed`

- `-1 < +1`

- `x22 > x23 // unsigned`

- `+4,294,967,295 > +1`

# Branching on Conditions other than (Not) Equal

- **Set Less Than (slt)**
  - `slt dst, reg1, reg2`
  - If value in `reg1` < value in `reg2`, `dst` = 1, else 0
- **Set Less Than Immediate (slti)**
  - `slti dst, reg1, imm`
  - If value in `reg1` < `imm`, `dst` = 1, else 0

# Recommended Reading

- RISC-V Edition - Computer Organization and Design\_ The Hardware Software Interface - David A. Patterson, John L. Hennessy:
  - Chapter-2