

# EE-421: Digital Systems Design

## Pipelining

Dr. Rehan Ahmed [rehan.ahmed@seecs.edu.pk]

# What is Pipelining?

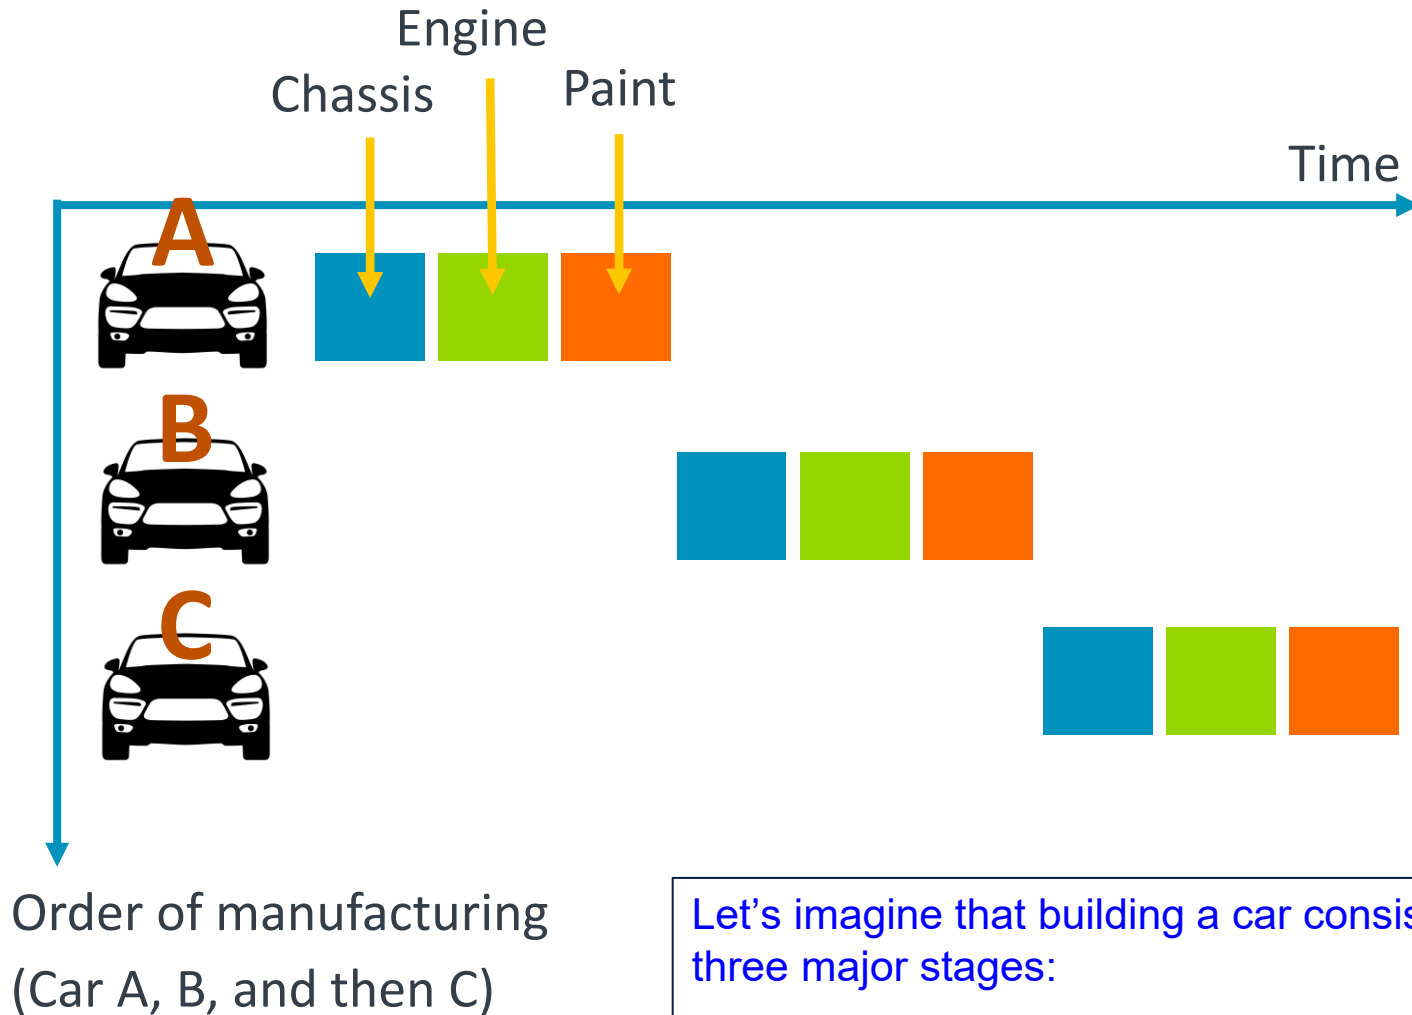
- Basic Idea:
  - We arrange for the different phases of execution to be overlapped.
  - We aim to exploit “temporal” parallelism.
- How are latency and throughput affected?
  - Let’s understand through examples (next)

# Pipelining: Assembly Line Example



Volkswagen Beetle Assembly Line  
(By [Alden Jewell](#), license: [CC BY 2.0](#))

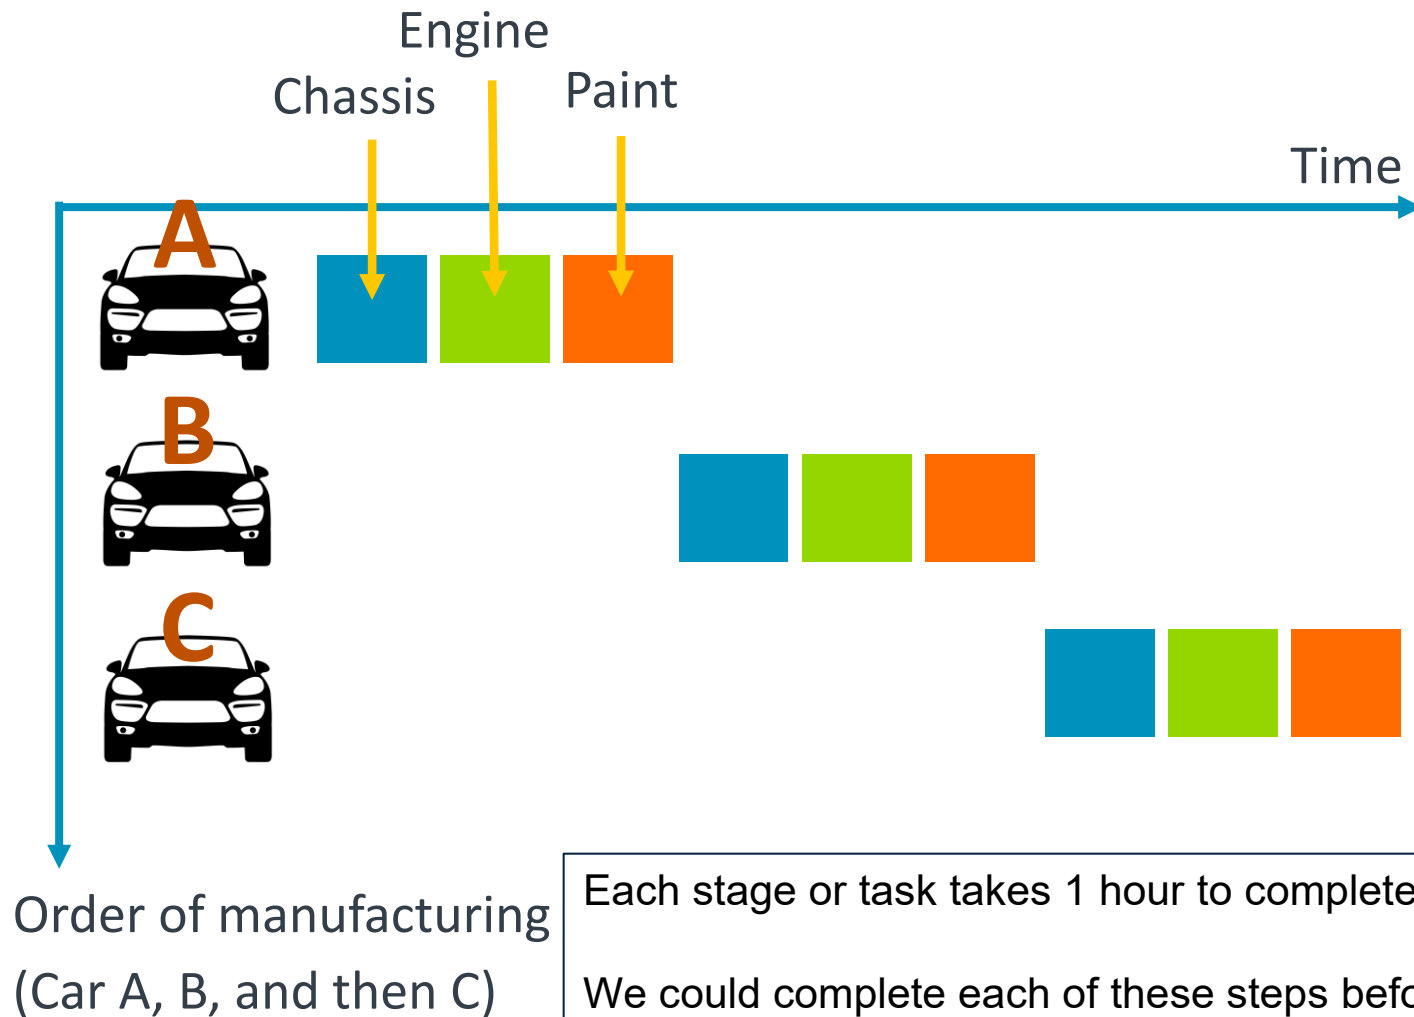
# Sequential Car Manufacturing



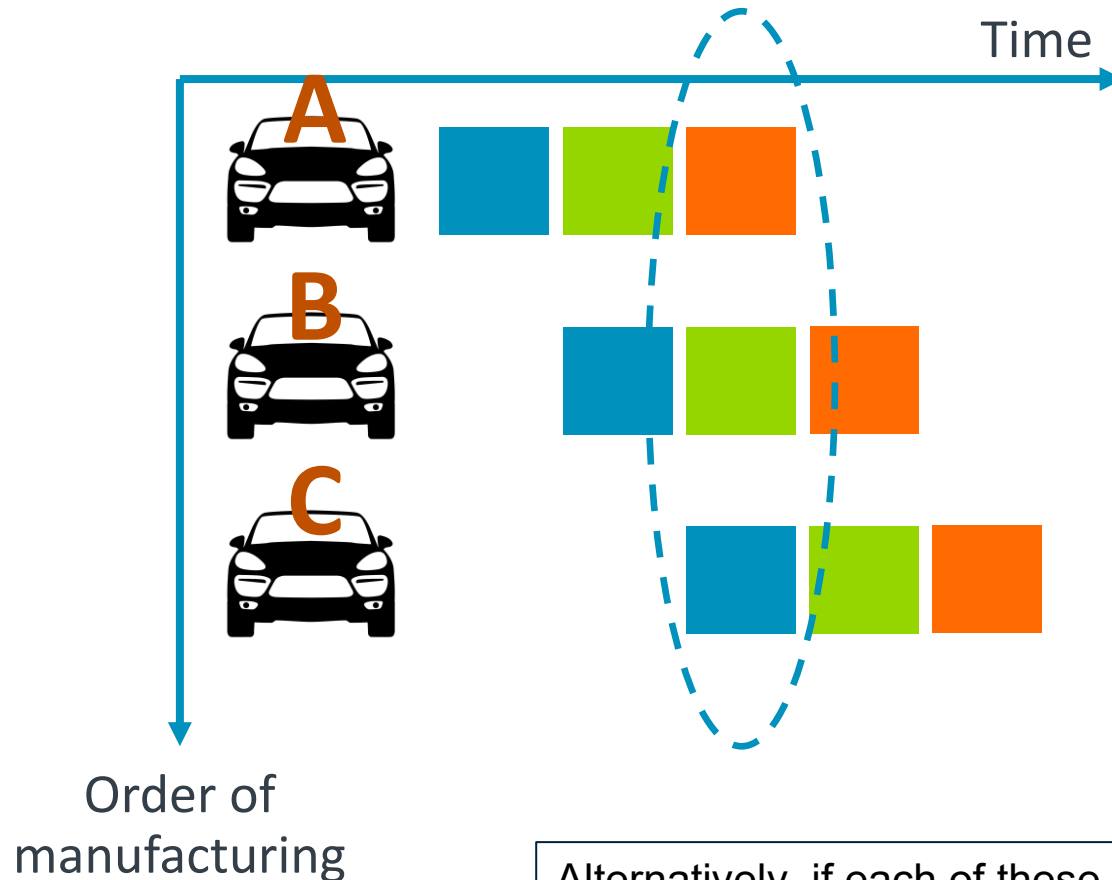
Let's imagine that building a car consists of three major stages:

- (1) constructing the chassis
- (2) installing the engine
- (3) painting.

# Sequential Car Manufacturing



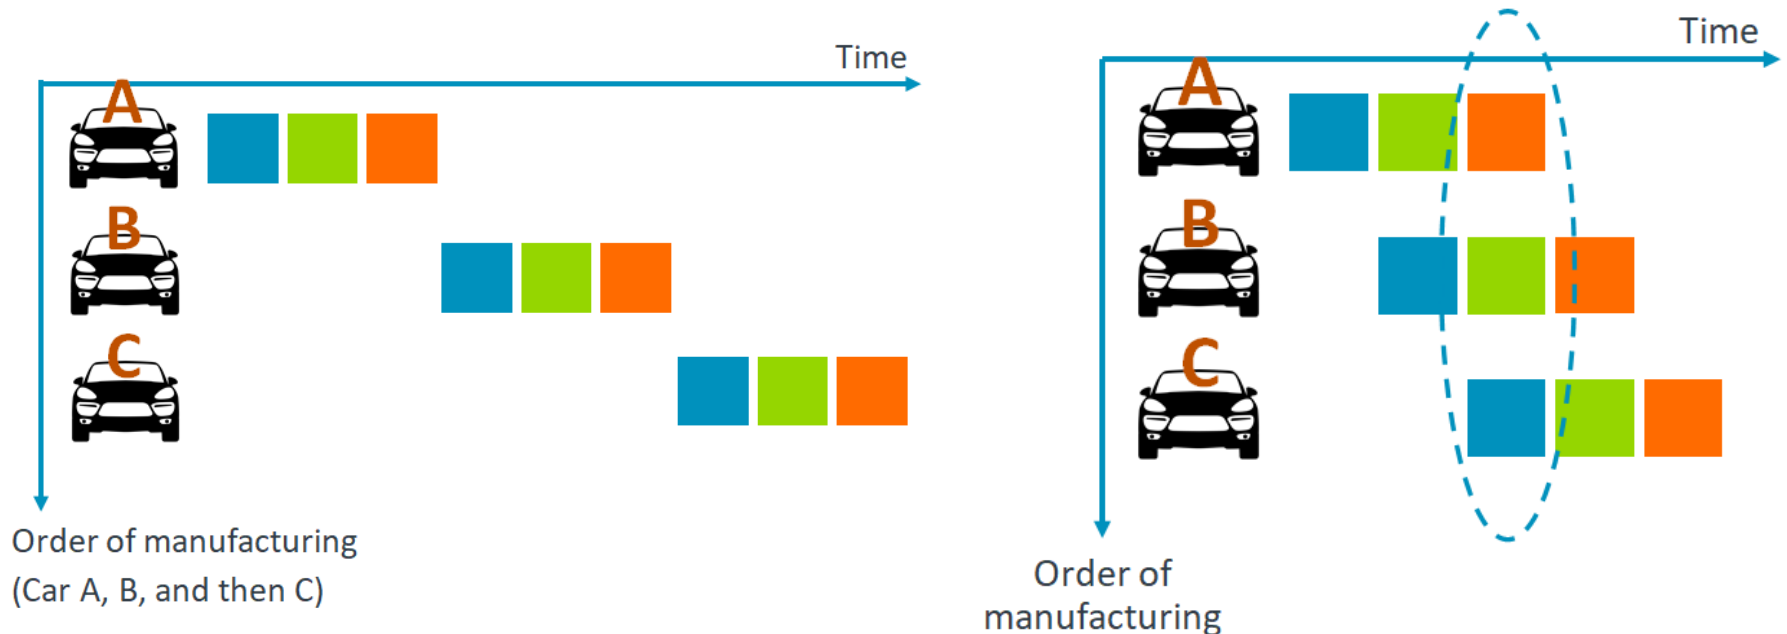
# Pipelined Car Manufacturing



Alternatively, if each of these tasks is completed at a different station, we could pipeline the cars construction. Now, the construction of a new car is started as soon as the first task is complete. Once our pipeline is full, a car will be completed every hour. Our throughput is now 24 cars per day.



# Pipelining Lessons

- Of course, a single car is not produced more quickly, **latency** may actually **increase** (due to unbalanced stages).
- Nevertheless, the rate at which cars are produced, our **throughput**, is much **higher** than if we constructed a single car at a time.



# Pipelining: Laundry Example




D E N S



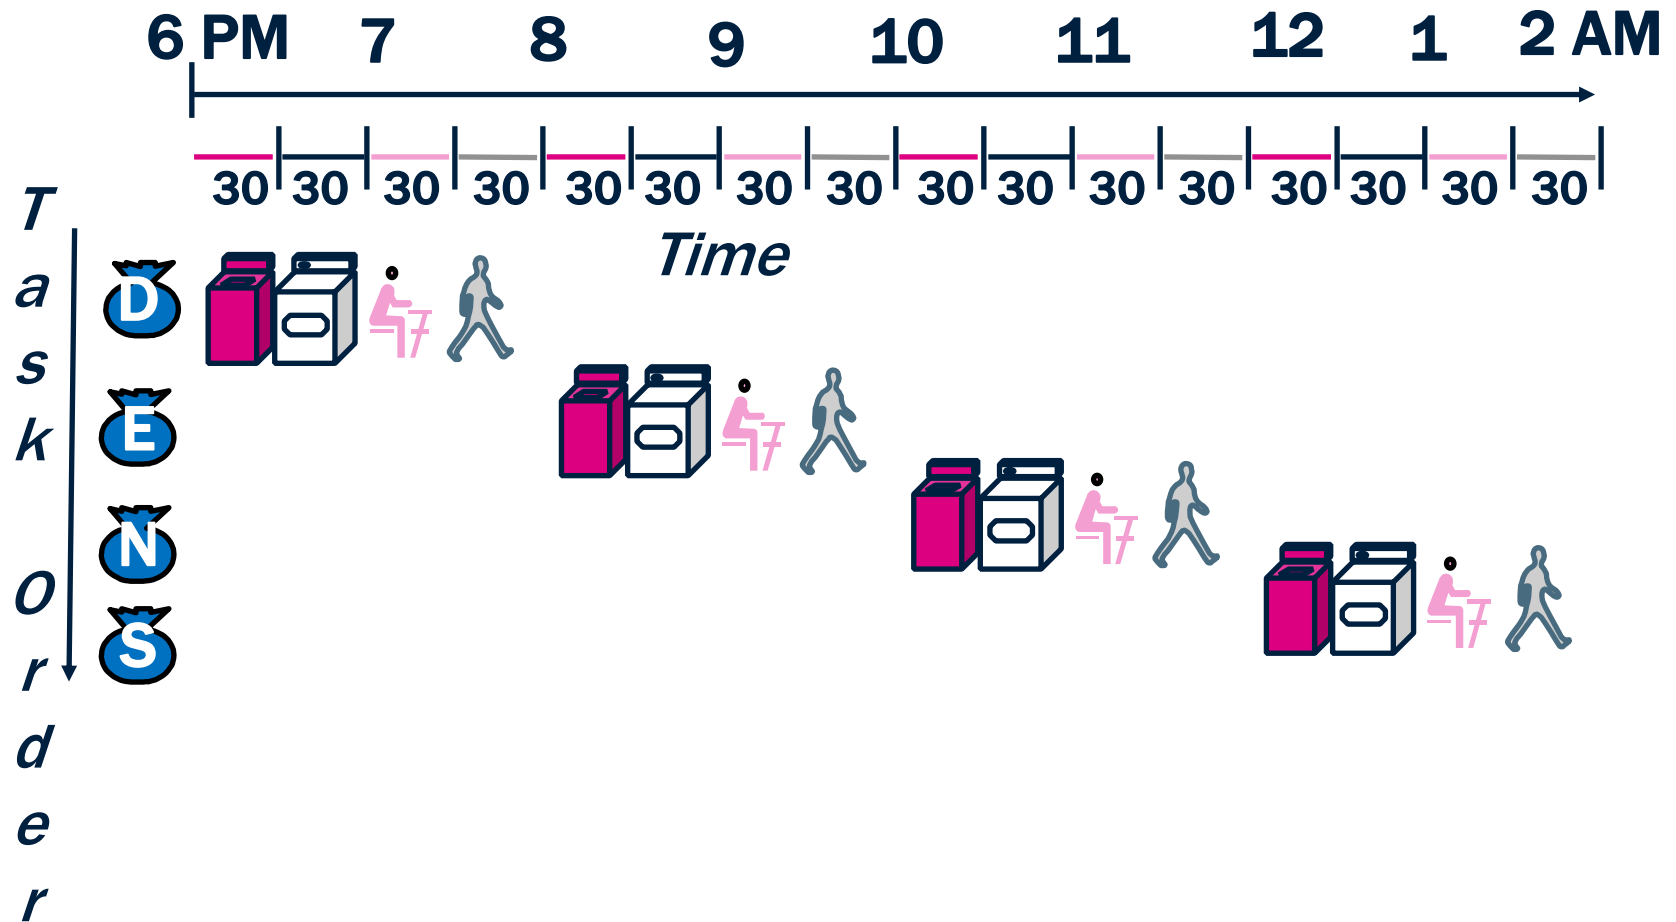


# Pipeline Analogy: Doing Laundry

- Damon, Emaan, Nick, and Steven each have one load of clothes to wash, dry, fold, and put away
  - Washer takes 30 minutes
  - Dryer takes 30 minutes
  - “Folder” takes 30 minutes
  - “Stasher” takes 30 minutes to put clothes into drawers

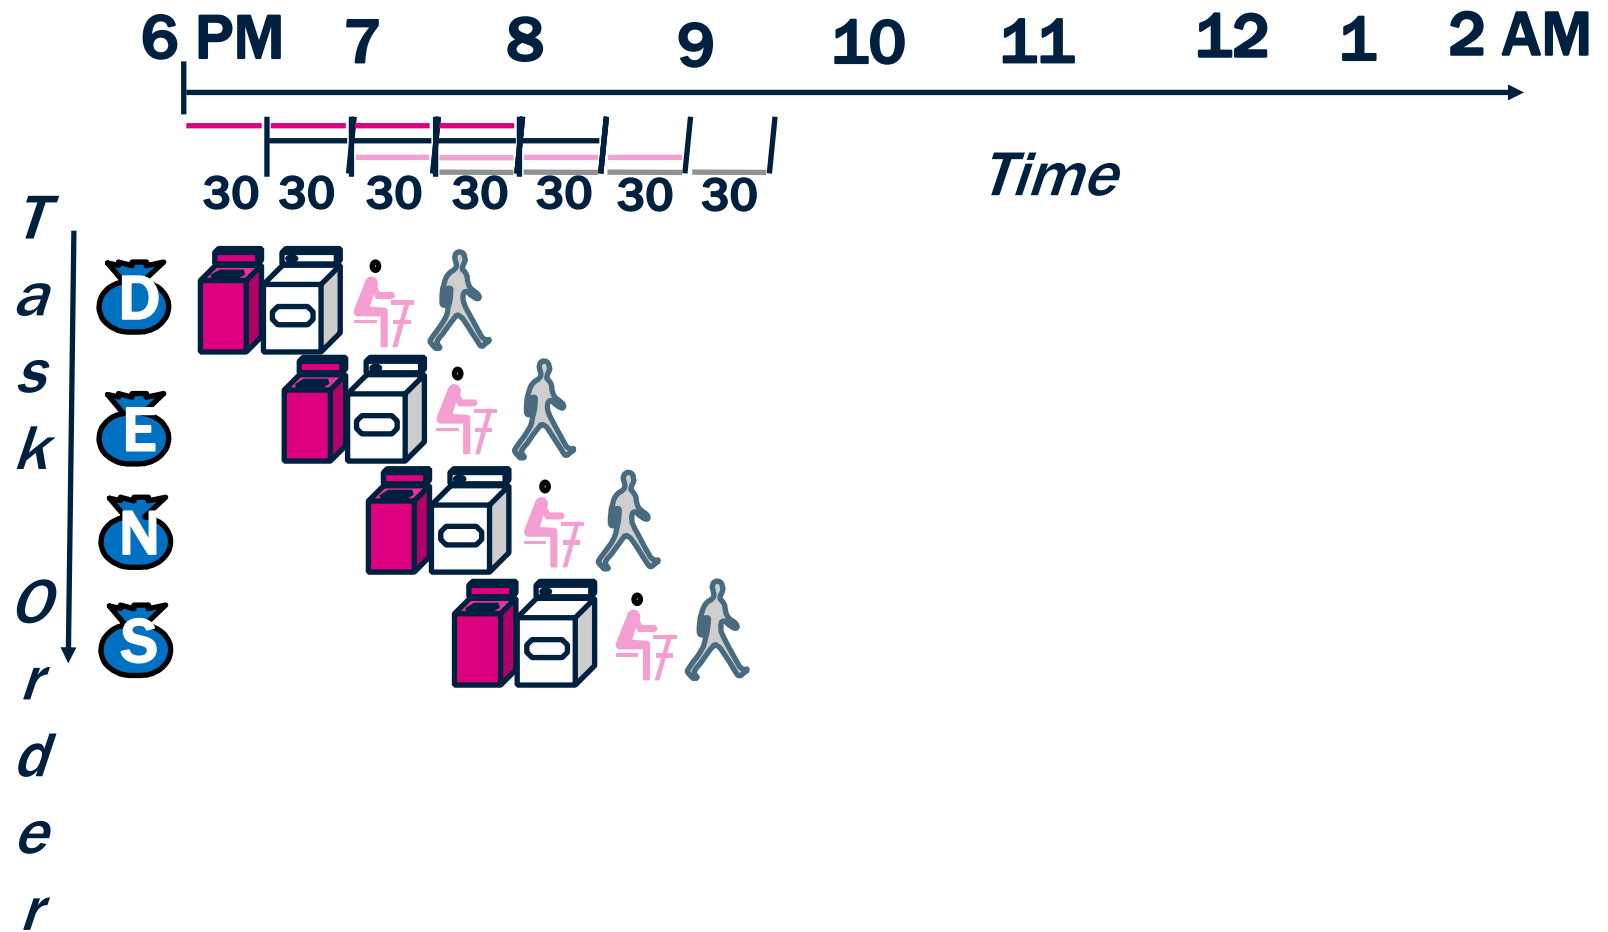


# Sequential Laundry



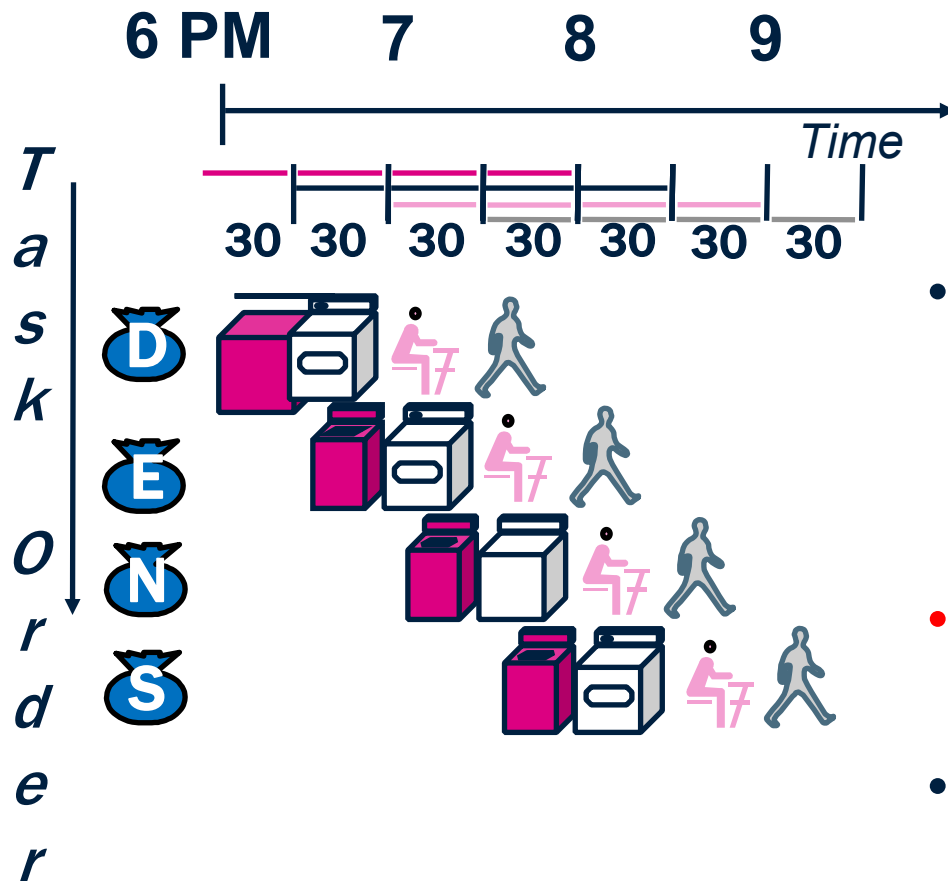
Sequential laundry takes 8 hours for 4 loads

# Pipelined Laundry



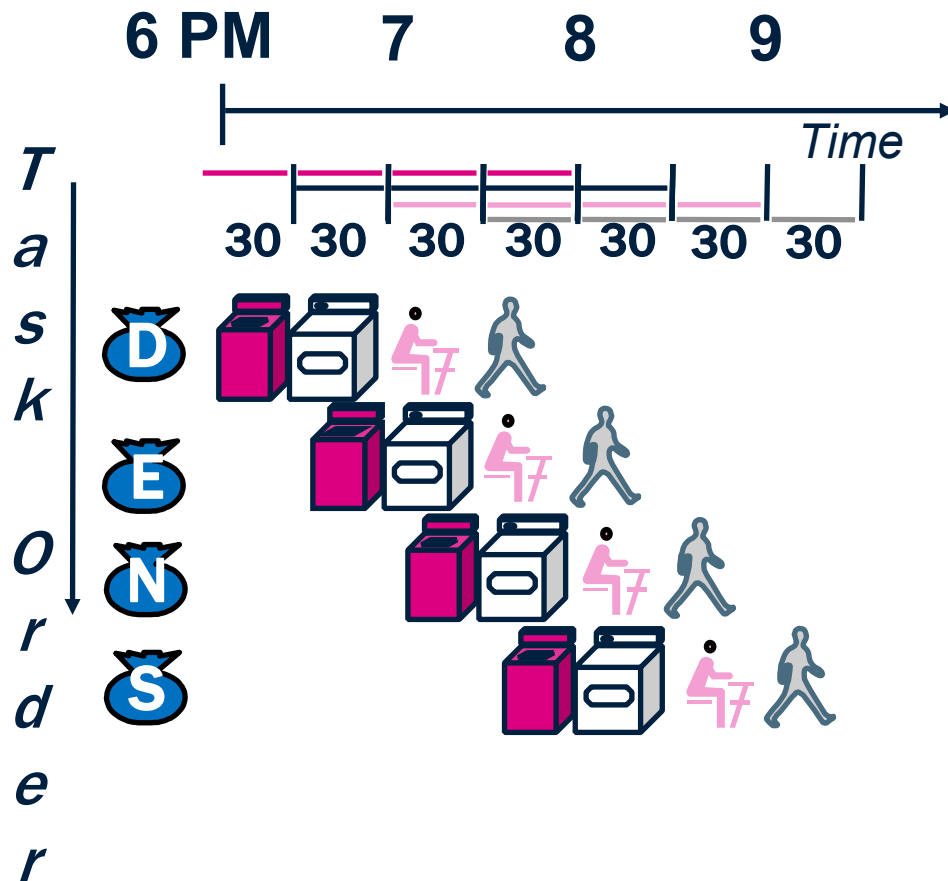
*Pipelined laundry takes 3.5 hours for 4 loads!*

# Pipelining Lessons (1/2)



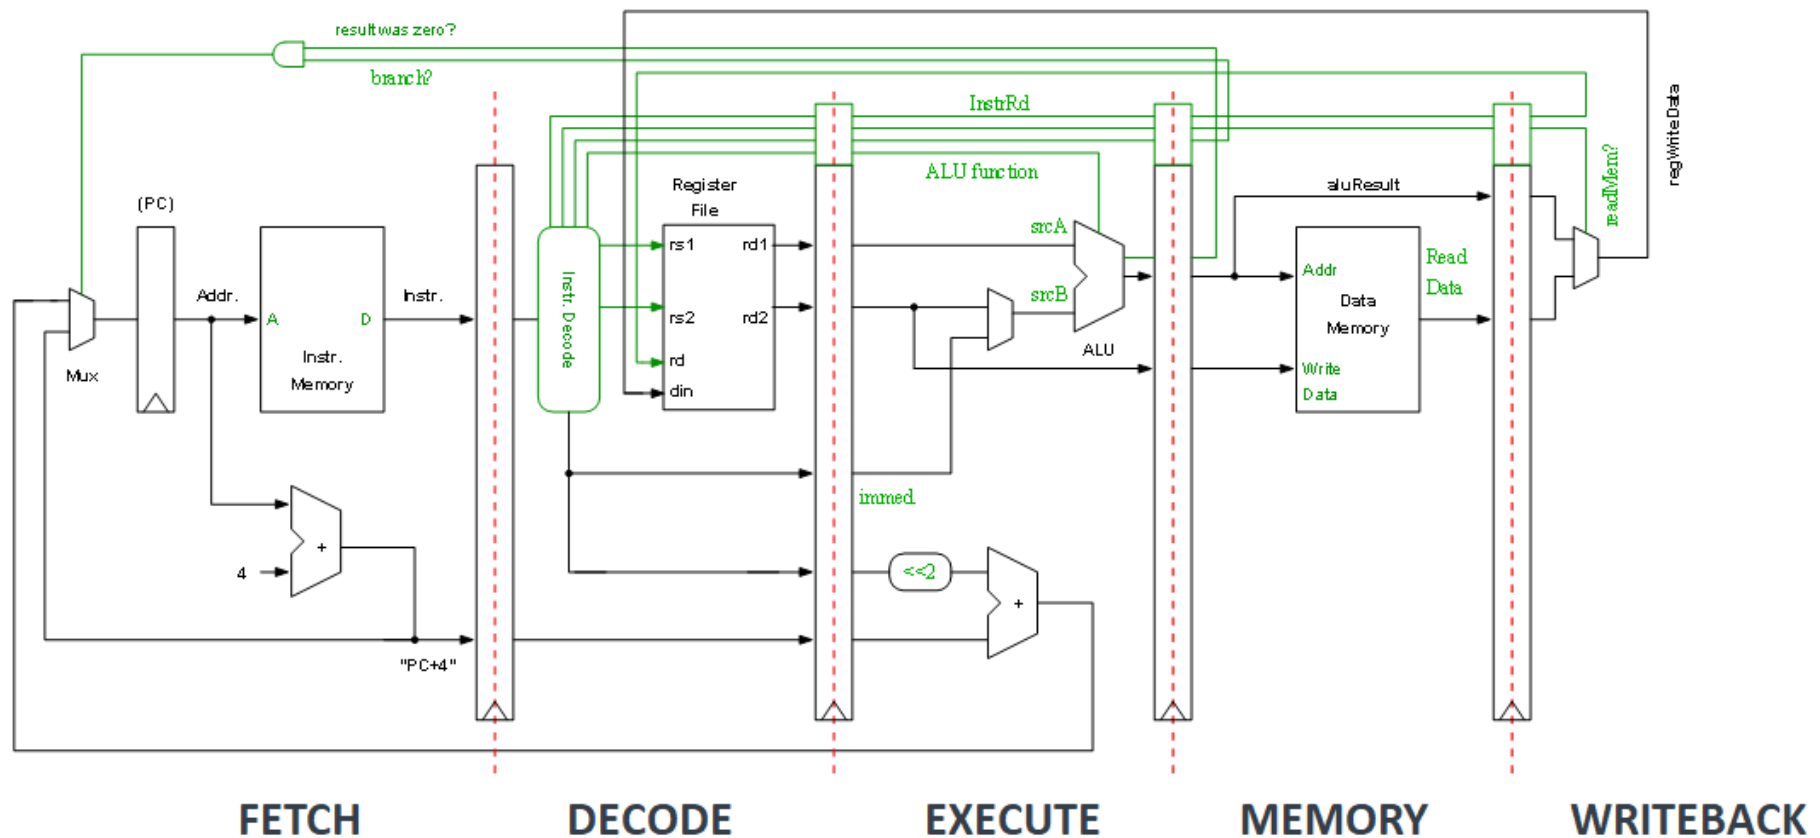
- Pipelining **doesn't help latency of single task**, just **throughput of entire workload is increased!**
- *Multiple* tasks operating simultaneously using different resources
- **Potential speedup = number of pipeline stages**
- Speedup reduced by time to **fill** and **drain** the pipeline:  
8 hours/3.5 hours or 2.3X  
v. potential 4X in this example

# Pipelining Lessons (2/2)

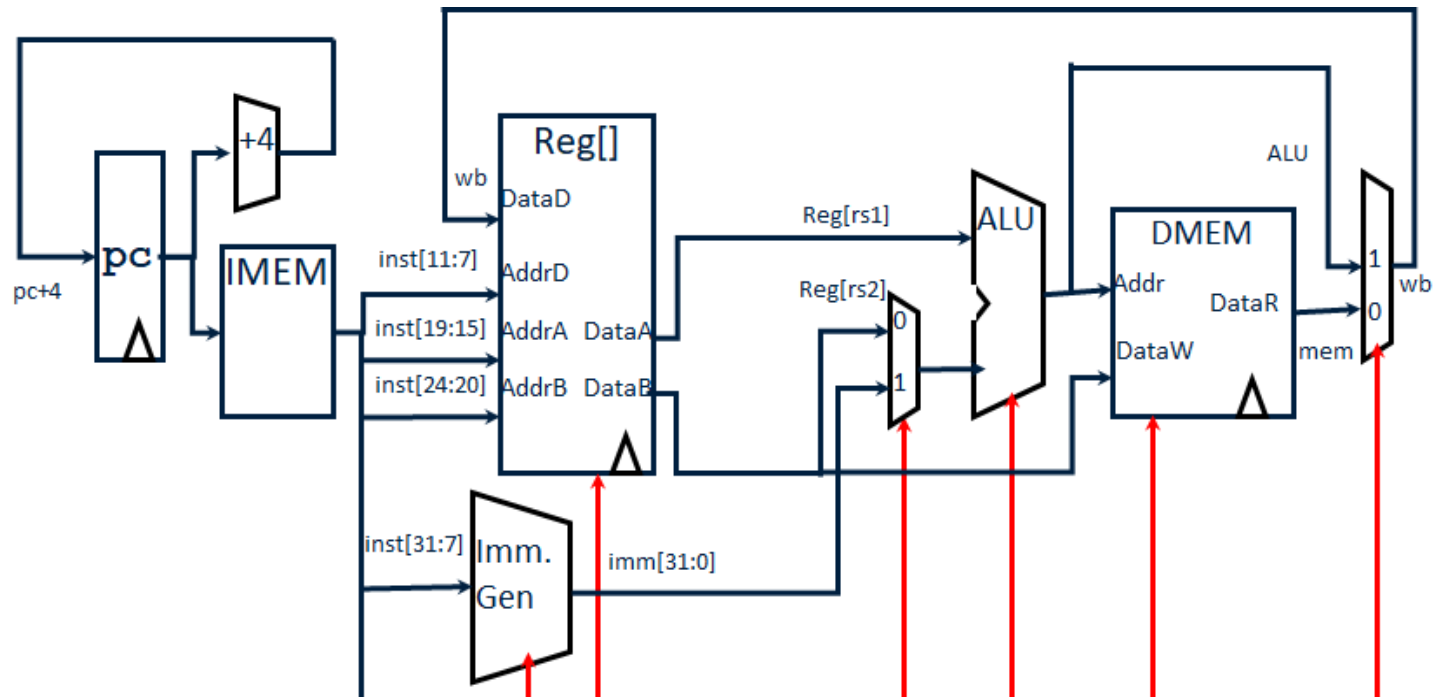


- Suppose new Washer takes 20 minutes, new Stasher takes 20 minutes. How much faster is pipeline?
  - Pipeline rate limited by *slowest* pipeline stage
  - Unbalanced lengths of pipeline stages reduces speedup

# Pipelining: Processor Datapath Example



# Single-Cycle Datapath



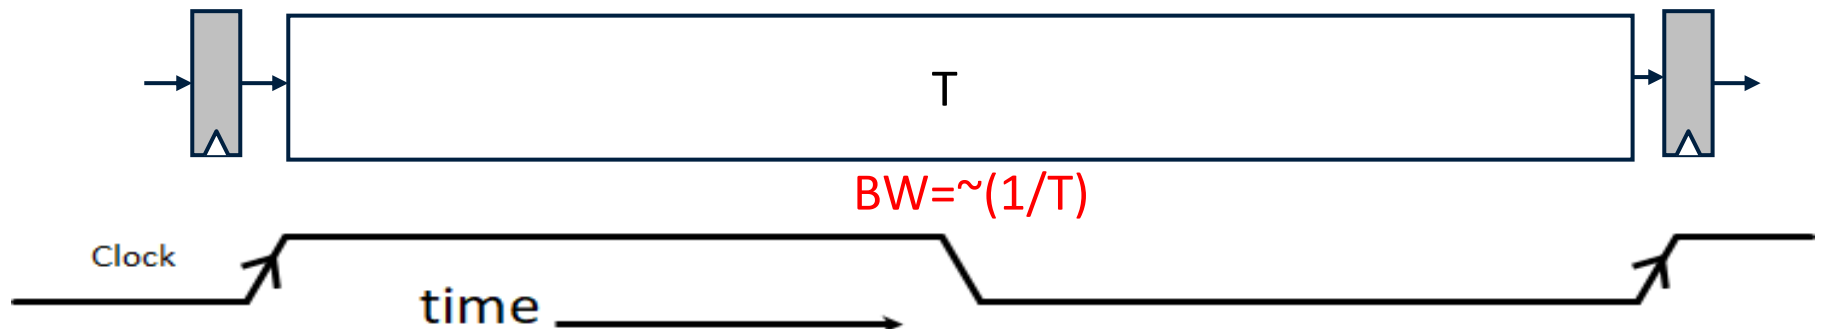
**FETCH**

**DECODE**

**EXECUTE**

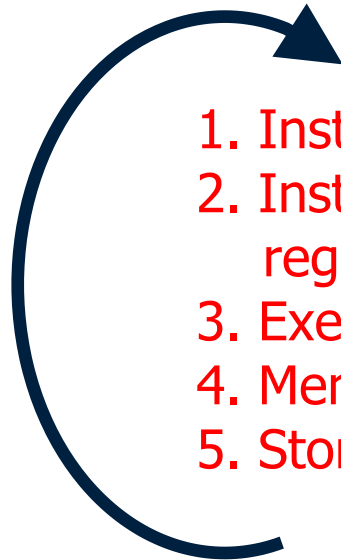
**MEMORY**

**WRITEBACK**



# Single-Cycle Datapath

- Our single-cycle processor executes each instruction in one clock cycle, i.e., it has a **Clocks Per Instruction (CPI)** of 1.

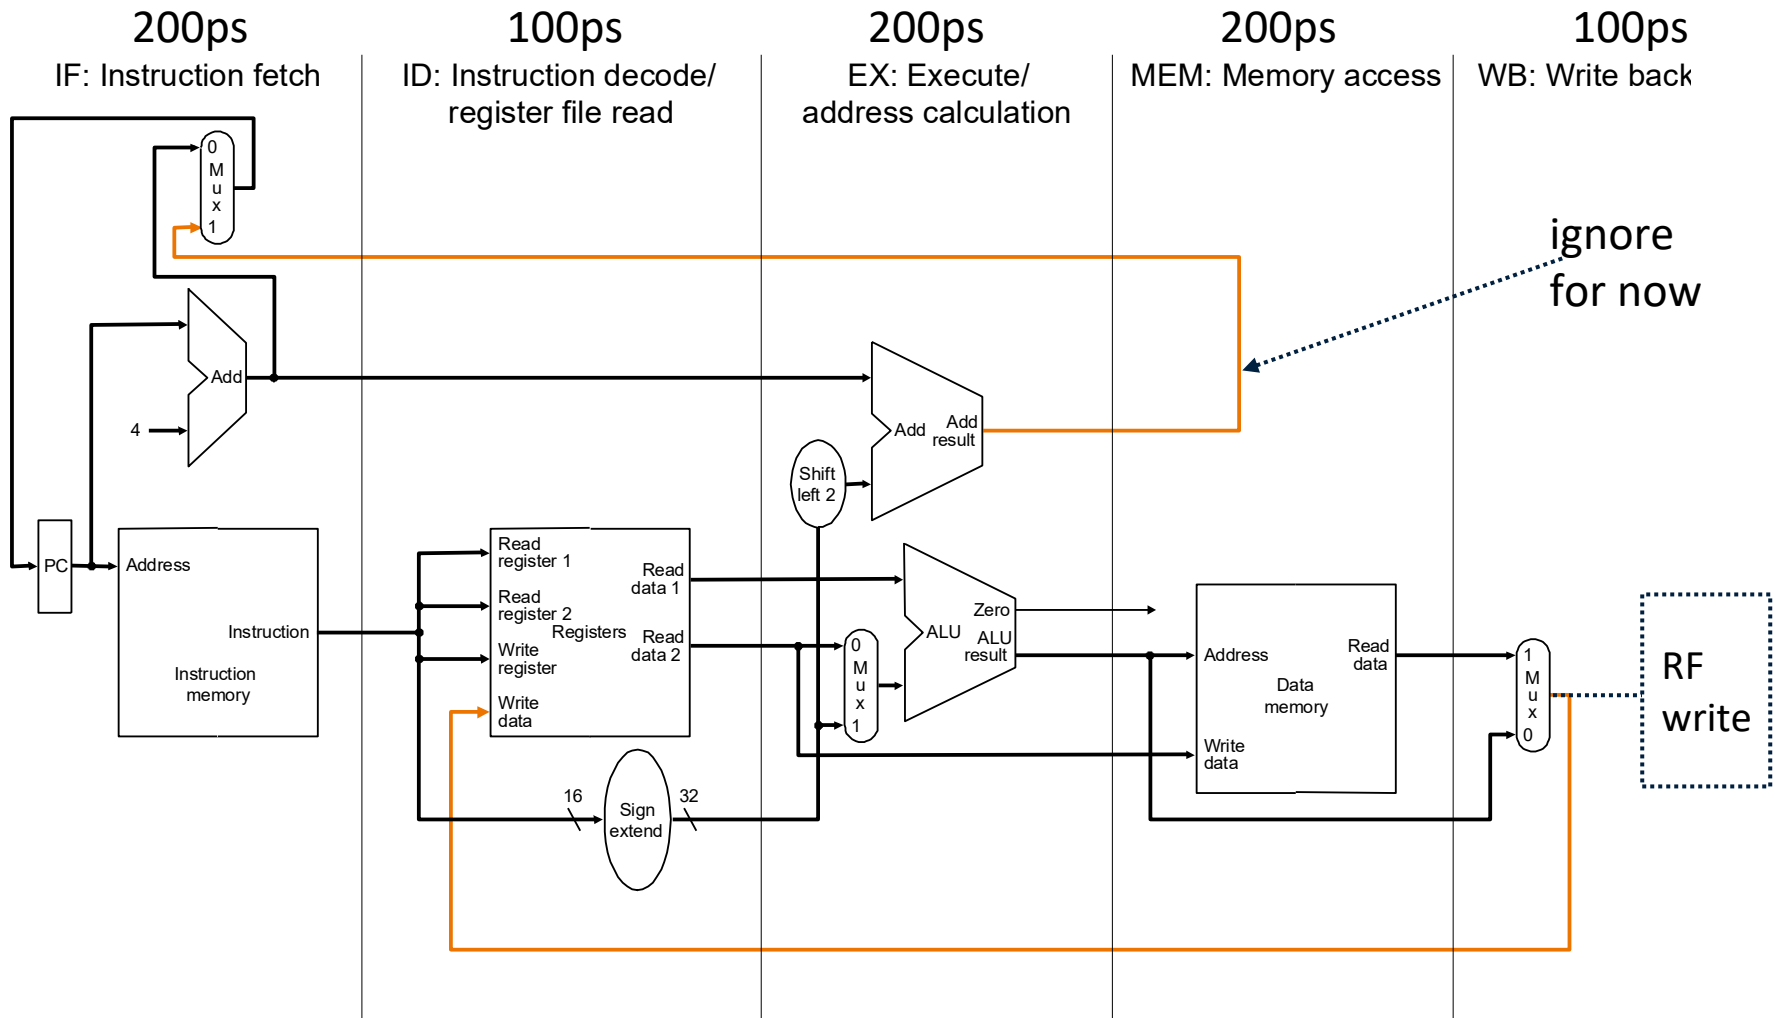


1. Instruction fetch (IF)
2. Instruction decode and register operand fetch (ID/RF)
3. Execute/Evaluate memory address (EX/AG)
4. Memory operand fetch (MEM)
5. Store/writeback result (WB)






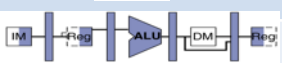
- How might we **improve** our **clock frequency** **without** significantly **increasing CPI**?
  - We can break the execution of instructions into stages and overlap the execution of different instructions.
  - We need to ensure that the results produced by our new pipelined processor are no different to the unpipelined one.



# Pipelined Datapath: Dividing Into Stages



# Pipelined Datapath

Phase	Pictogram	$t_{step}$ Serial	$t_{cycle}$ Pipelined
Instruction Fetch		200 ps	200 ps
Reg Read		100 ps	200 ps
ALU		200 ps	200 ps
Memory		200 ps	200 ps
Register Write		100 ps	200 ps
$t_{instruction}$		800 ps	1000 ps

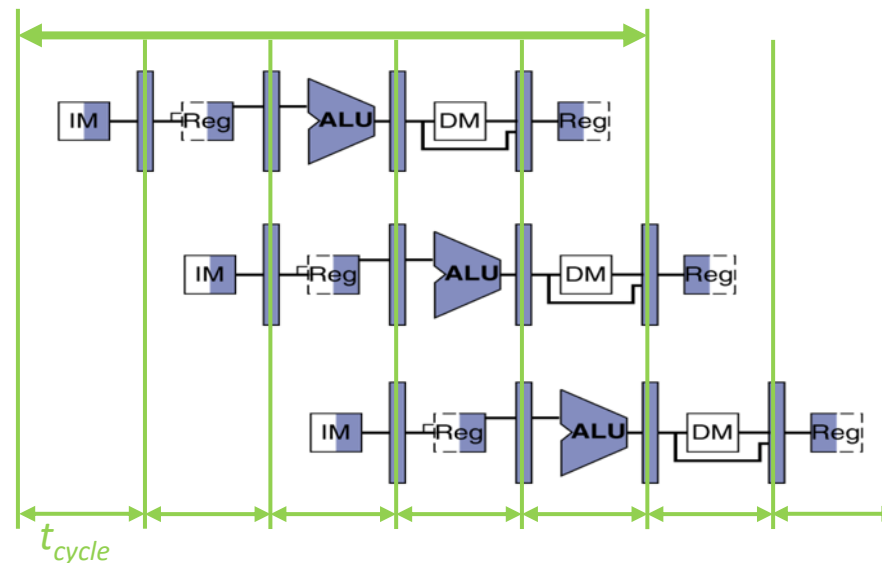
instruction sequence



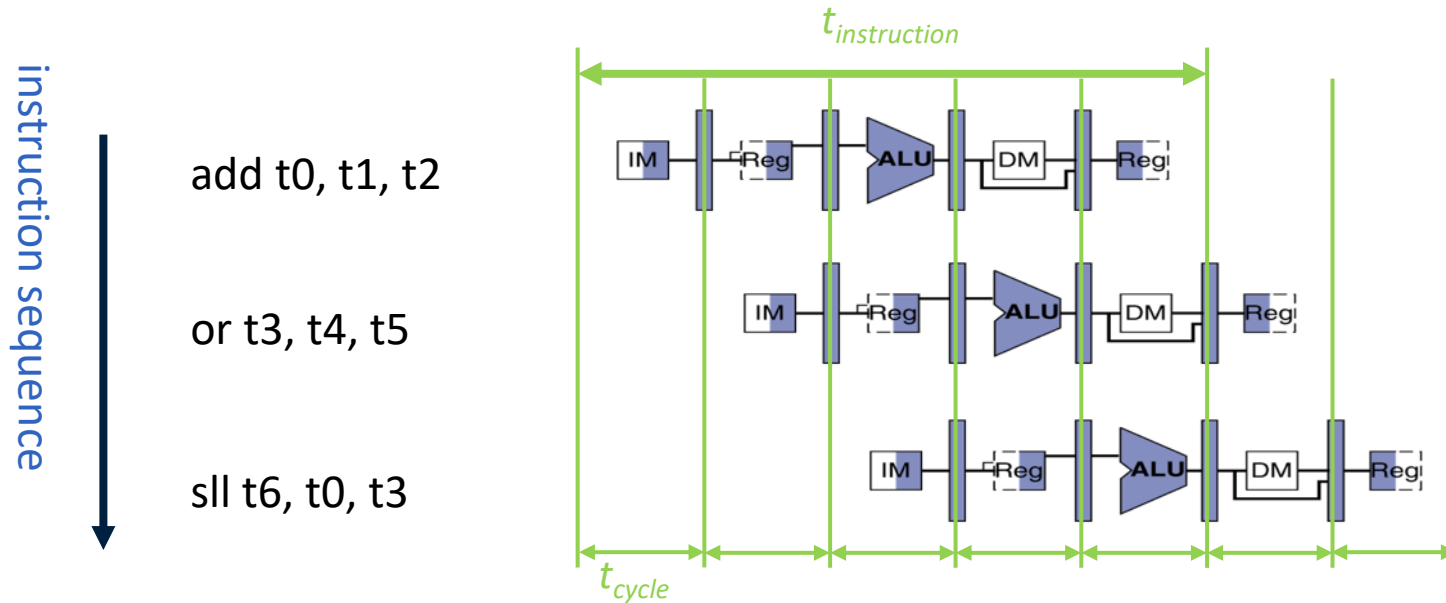
add t0, t1, t2

or t3, t4, t5

sll t6, t0, t3



# Pipelined Datapath

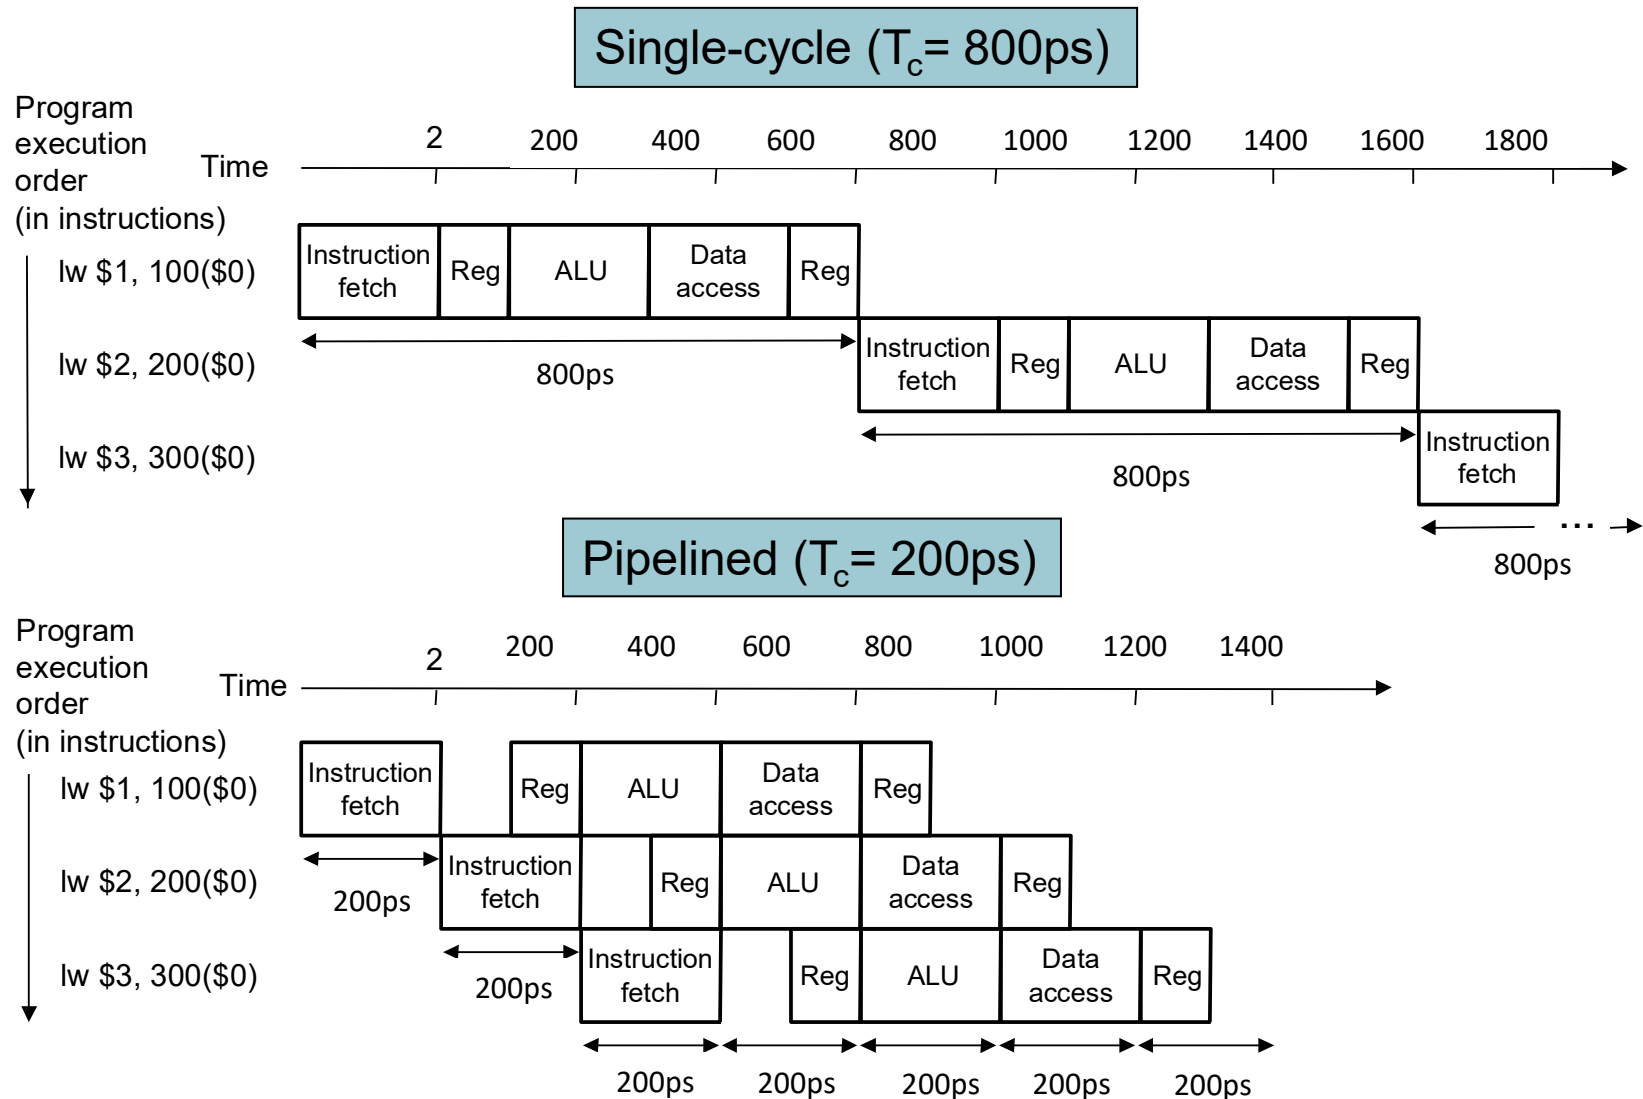


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, $f_s$	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x

# Pipelining Lessons

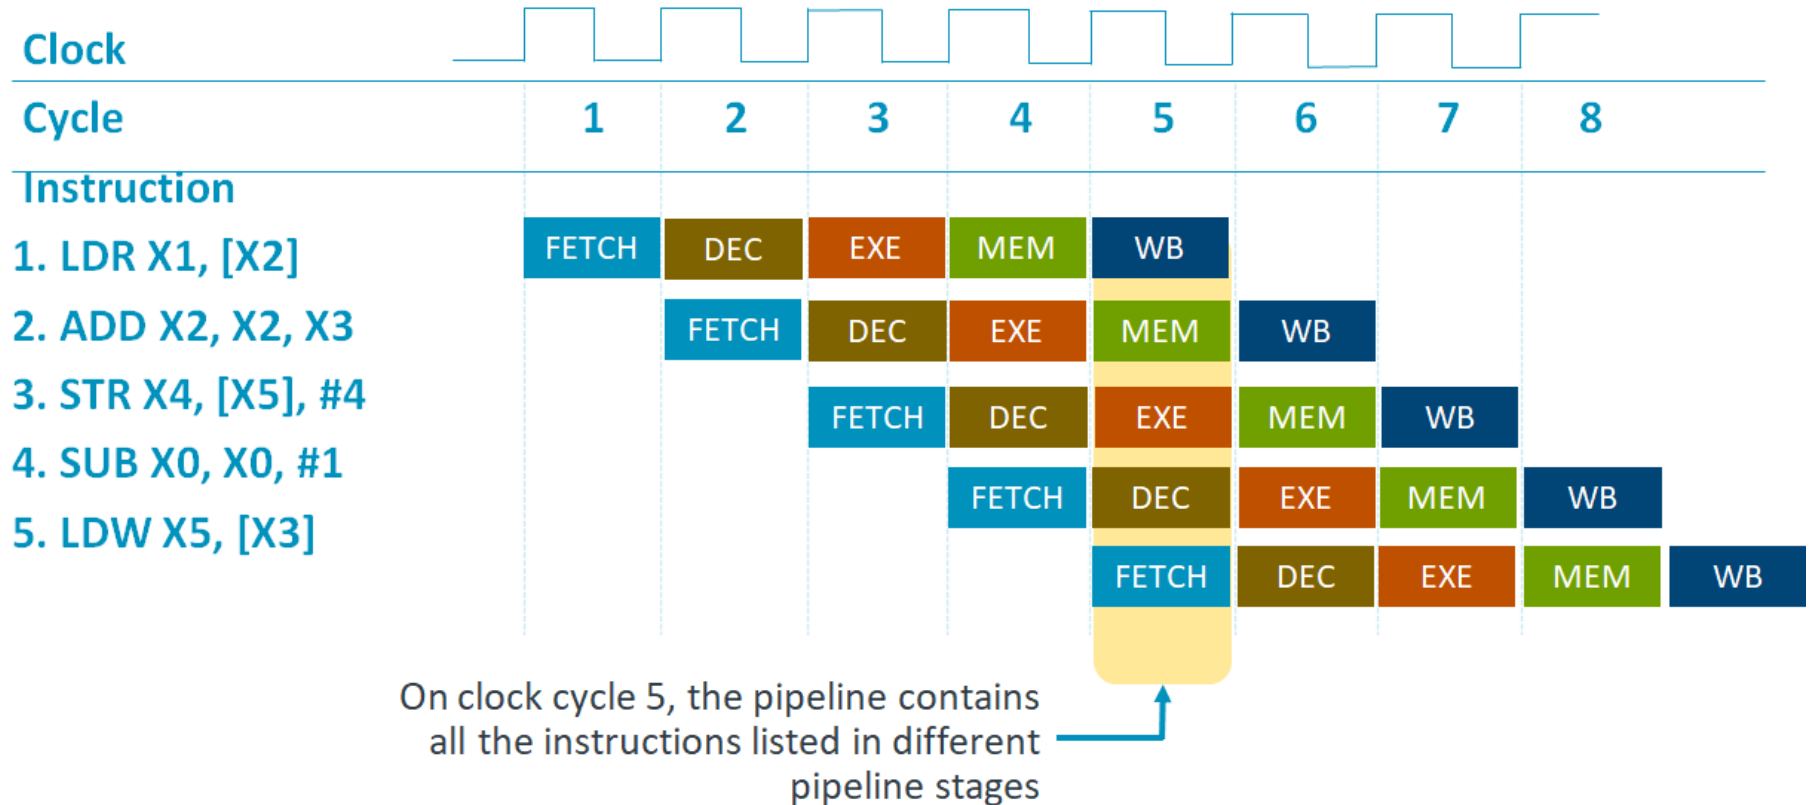
- Use  $T_c$  (“time between completion of instructions”) to measure speedup
  - $T_{c, \text{pipelined}} \geq \frac{T_{c, \text{single-cycle}}}{\text{Number of stages}}$
  - Equality only achieved if stages are *balanced* (i.e. take the same amount of time)
- If not balanced, speedup is reduced
- Speedup due to increased *throughput*
  - *Latency* for each instruction does not decrease

# Instruction Pipeline Throughput



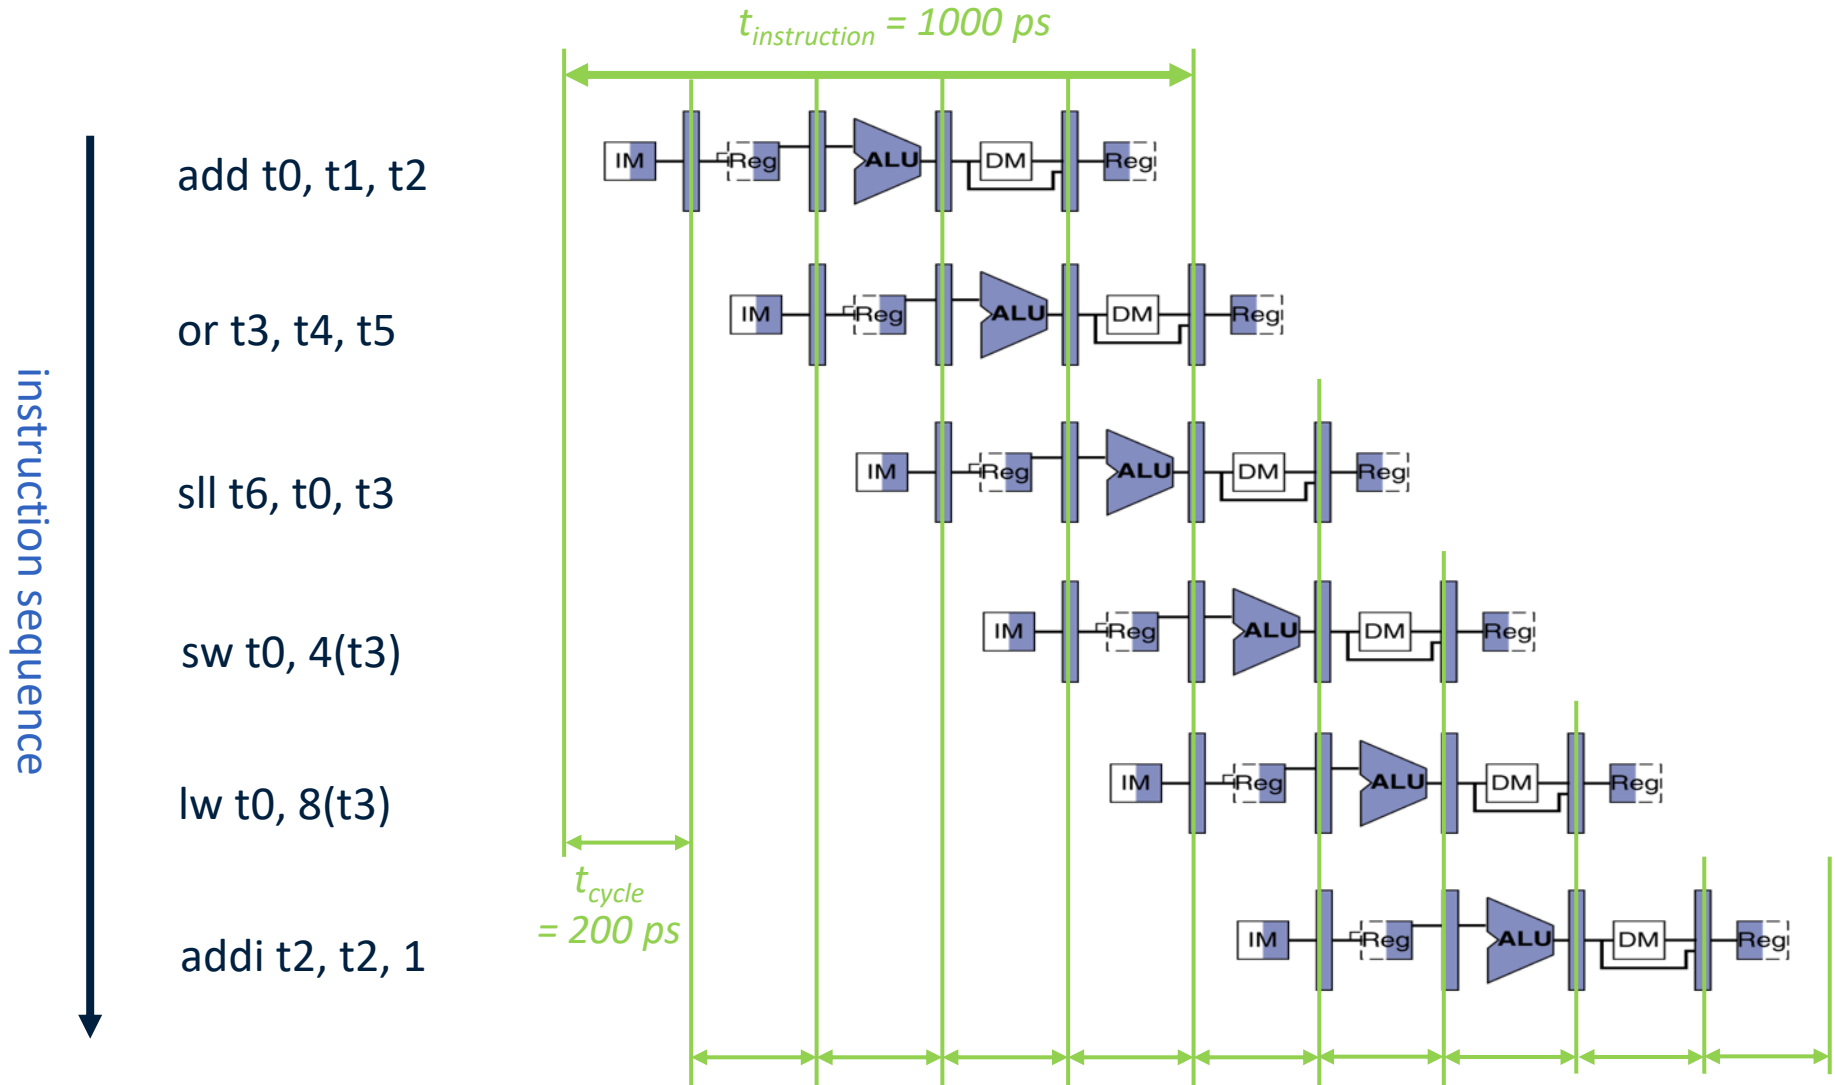
5-stage speedup is 4, not 5 as predicted by the ideal model. Why?

# The Pipeline in Action






# Sequential vs Simultaneous

*What happens sequentially, what happens simultaneously?*



# Implementing Pipelining





# Summary

- Pipelining improves performance by *increasing instruction throughput, in contrast to decreasing the execution time of an individual instruction*, but instruction throughput is the important metric because real programs execute billions of instructions.

# Recommended Reading

- Computer Organization and Design\_ The Hardware Software Interface - David A. Patterson, John L. Hennessy:
  - Chapter-4:
    - 4.5
    - 4.6 -> Graphically representing pipeline

# THANK YOU

